

Dissecting the NVidia Turing T4 GPU via Microbenchmarking

Technical Report

Zhe Jia
Marco Maggioni
Jeffrey Smith
Daniele Paolo Scarpazza



High Performance Computing R&D Team
Citadel, 131 S. Dearborn St., Chicago, Ill. 60603.

Copyright © 2019, Citadel Enterprise Americas, LLC. All rights reserved.

Chicago, United States of America.

Citadel grants arXiv.org a perpetual, non-exclusive license to distribute this manuscript.

This presentation solely reflects the analyses and views of the authors. No recipient should interpret this presentation to represent the general views of Citadel or its personnel. Facts, analyses, and views presented herein have not been reviewed by, and may not reflect information known to other Citadel professionals.

This report strives to be objective, neutral and impartial. The authors did not receive any form of compensation or reimbursement in conjunction with this work, and are not employed by NVidia. All results in this report are meant to be reproducible by any researcher who recreates the experimental setup described by the authors, except for good-faith mistakes. No such mistakes are known to the authors at the time of publication. Findings are solely based only on the authors' experiments and understanding, independent of NVidia's confirmation of methods or data.

The authors welcome all feedback. The corresponding author is Daniele P. Scarpazza. Send feedback preferably to his e-mail address: daniele.scarpazza@citadel.com.

Edition	Date
First	March 19, 2019

The authors make this edition, and will make future editions, of this document available on the [arXiv.org](https://arxiv.org) e-print service owned and operated by Cornell University. The authors presented material contained in this report at the 2019 [GPU Technology Conference](#), March 18-21, 2019, San Jose, California.

All product names, trademarks and registered trademarks are property of their respective owners. All company, product and service names mentioned are for identification purposes only. Use of these names, trademarks and brands does not imply endorsement.

Summary

In 2019, the rapid rate at which GPU manufacturers refresh their designs, coupled with their reluctance to disclose microarchitectural details, is still a hurdle for those software designers who want to extract the highest possible performance from GPUs.

Last year, these very reasons motivated us to dissect the Volta GPU architecture using microbenchmarks. We presented our findings at NVidia’s GPU Technology Conference (GTC2018) [1] and published them in a technical report [2].

The introduction in August 2018 of Turing [3], NVidia’s latest architecture, pressed us to update our study. In this report, we examine Turing and compare it quantitatively against previous NVidia GPU generations. Specifically, we study the T4 GPU: a low-power, small form-factor board aiming at inference applications. We describe its improvements against its inference-oriented predecessor: the P4 GPU based on the Pascal architecture. Both T4 and P4 GPUs achieve significantly higher frequency-per-Watt figures than their full-size counterparts.

We study the performance of the T4’s Tensor Cores, finding a much higher throughput on low-precision operands than on the P4 GPU. We reveal that Turing introduces new instructions that express matrix math more succinctly. We map Turing’s instruction space, finding the same encoding as Volta, and additional instructions. We reveal that the Turing TU104 chip has the same memory hierarchy depth as the Volta GV100; cache levels sizes on the TU104 are frequently twice as large as those found on the Pascal GP104. We benchmark each constituent of the T4 memory hierarchy and find substantial overall performance improvements over its P4 predecessor. We studied how clock throttling affects compute-intensive workloads that hit power or thermal limits.

Many of our findings are novel, published here for the first time. All of them can guide high-performance software developers get closer to the GPU’s peak performance, as we illustrate with examples.

Chapter 1

Low-level details make a difference

In this section, we use a practical example to motivate our claim that a deep understanding of the architecture can help developers achieve substantial speed-ups in certain, specific scenarios, although at the cost of significant development effort.

It takes disproportionate effort to optimize code on the basis of a deep understanding of its target architecture. This approach frequently resorts to writing inline PTX assembly and, when pushed to its extreme, to patching binary code in the pursuit of specific SASS assembly that the compiler won't emit, following undocumented instruction encoding formats, without any support from NVidia's toolchain.

Whether the gains are worth the effort is a central question but, ultimately, one that only you can answer depending on your unique circumstances and pressure for performance. For a large majority of GPU software developers, *the answer is no*:

- if one of the mature libraries provided by NVidia (such as cuBlas, cuFFT, cuSparse, cuRand, cuDNN, etc.) covers the computation desired, the performance obtained is close to ideal in most circumstances;
- in other applications, for which developers write CUDA code, NVCC usually emits efficient machine code, if the source code is written sufficiently well.

Rare are the cases where the pressure for performance justifies extreme, low-level optimization. A prolific line of research has traditionally focused on understanding GPU instruction encoding [4, 5, 6] precisely to improve the performance of compute kernels [7, 8, 9]. Our prior work on Volta [2] also offered such an example: we patched compiler-emitted code so that it used the register cache better, and achieved a 15% higher floating-point arithmetic throughput.




Figure 1.1: Performance of our improved `saxpy` implementation that uses 128-bit wide global memory access instructions, compared with `cublasSaxpy` from NVidia’s cuBlas library, that uses 32- and 64-bit wide instructions. Elapsed time in microseconds; lower is better.

This time, we show that the knowledge of Turing’s instructions allows designers to improve the performance of a common linear algebra function (i.e., BLAS `?axpy`), whose library implementation for single-precision operands contains memory access operations limited to a 64-bit width. We show a simple replacement that uses vectorized 128-bit accesses, and improves performance substantially (see Fig. 1.1).

The `?axpy` workload we chose performs a scaled, element-wise vector-vector sum, i.e.:

$$\vec{y} := a \cdot \vec{x} + \vec{y}$$

where a is the scale factor. The workload is quite obviously of low arithmetic intensity, which makes it memory bound. This means that any efficiency gain in its access to global memory translates into a direct overall speedup.

The most straightforward way to compute this workload in C on a CPU code is invoking BLAS function `cblas_?axpy` (where ‘?’ is a placeholder for the operand type, i.e., `s`, `d`, `c`, or `z`, respectively, for single- or double-precision, real or complex operands). In CUDA, the corresponding function

is `cublas?axpy`, from NVidia's cuBlas library. For simplicity, we only focus on its single-precision variant `cublasSaxpy`.

In our analysis, function `cublasSaxpy` calls, in turn, `axpy_kernel_val` which, in its single-precisions implementations shipped with the CUDA SDK version 10.0, contain load and store instructions (from/to global memory) no wider than 64 bits. An excerpt of one such implementation follows:

```
Function : void axpy_kernel_val<float2, float2, 0>(
    cublasAxyParamsVal<float2, float2, float2>)
    .headerflags      @"EF_CUDA_SM75 EF_CUDA_PTX_SM(EF_CUDA_SM75) "
...
    /*01f0*/          LDG.E.64.SYS R2, [R2] ;
    /*0200*/          LDG.E.64.SYS R8, [R4] ;
...
    /*0330*/          LDG.E.64.SYS R2, [R2] ;
    /*0340*/          LDG.E.64.SYS R8, [R4] ;
...
```

This access width is suboptimal on Turing, especially considering that:

1. Turing has only half as many load/store units per SM as Volta (16 vs. 32, as per public NVidia information [10, 3]);
2. a T4 device supports only half as many threads per SM than a V100 device (1,024 vs. 2,048); it is therefore harder to saturate the available memory bandwidth on Turing by just increasing block count.

As a consequence, loading wider words per instruction is an effective strategy to increase memory access throughput. We do so in our improved implementation of the `Saxpy` kernel, that uses 128-bit vectorized memory access instructions:

```
__global__ void improved_Saxpy( float *d_y, const float *d_x,
    const float alpha, const uint32_t arraySize)
{
    // every thread process 4 elements at a time
    uint32_t tid = (threadIdx.x+blockIdx.x*blockDim.x)*4;
    // the elements that all threads on GPU can process at a time
    uint32_t dim = gridDim.x*blockDim.x*4;

    for(uint32_t i = tid; i < arraySize; i += dim)
        asm volatile ("{\t\n"
            // registers to store input operands
            ".reg .f32 a1,b1,c1,d1;\n\t"
            ".reg .f32 a2,b2,c2,d2;\n\t"

            // loading with vectorized, 128-bit instructions
            "ld.global.v4.f32 {a1,b1,c1,d1}, [%0];\n\t"
            "ld.global.v4.f32 {a2,b2,c2,d2}, [%1];\n\t"

            // core math operations
            "fma.rn.f32 a2,a1,%2,a2;\n\t"
            "fma.rn.f32 b2,b1,%2,b2;\n\t"
            "fma.rn.f32 c2,c1,%2,c2;\n\t"
            "fma.rn.f32 d2,d1,%2,d2;\n\t"

            // storing results with a vectorized, 128-bit write instruction
            "st.global.v4.f32 [%1], {a2,b2,c2,d2};\n\t"
```

```

    "}" :: "l"(d_x+i), "l"(d_y+i), "f"(alpha) : "memory"
    );
}

```

The PTX assembly that you see inlined in our code above visibly uses 128-bit wide load and store instructions `ld.global.v4.f32` and `st.global.v4.f32`, capable of transferring a vector of four single-precision floating-point values at a time. (For simplicity and brevity, our implementation neglects arrays whose size is not a multiple of 4.)

Inspection of the corresponding SASS code emitted by NVCC confirms that global memory instructions are 128 bits wide:

```

.headerflags      @"EF_CUDA_SM75 EF_CUDA_PTX_SM(EF_CUDA_SM75) "
...
/*00d0*/          LDG.E.128.SYS R8, [R8] ;
/*00e0*/          LDG.E.128.SYS R4, [R2] ;
...
/*0110*/          FFMA R4, R8, c[0x0][0x170], R4 ;
/*0120*/          FFMA R5, R9, c[0x0][0x170], R5 ;
/*0130*/          FFMA R6, R10, c[0x0][0x170], R6 ;
/*0140*/          FFMA R7, R11, c[0x0][0x170], R7 ;
/*0150*/          STG.E.128.SYS [R2], R4 ;
...

```

The performance of this `improved_Saxpy` code proves to be significantly higher than the performance of `cublasSaxpy` (see Figure 1.1) except for trivially small arrays (<20 Kib), and it asymptotically tends to be almost twice as fast for large ones.

In summary, with this example we demonstrated optimization opportunities that are only accessible to a software designer who possesses in-depth knowledge of Turing's instruction set and an architectural-level understanding of its performance behavior: these goals are the very subjects of this report.

Chapter 2

How Turing encodes instructions

By systematically disassembling machine code that we hand-crafted and that we sampled from representative CUDA libraries, using `cuobjdump` and `nvdiasm`, we discovered the instruction encoding formats adopted across the different GPU architectures. Turing adopts the same format as Volta, which differs from that of Pascal and Maxwell which, in turn, is different from Kepler's, as we detail in this chapter.

Turing and Volta use 128 bits to encode both an instruction and its associated scheduling control information¹. This is a substantial departure from previous NVidia GPU architectures that used one word per instruction (64-bit) to encode pure instruction information, plus a separate 64-bit word every few instructions, to encode control information associated those instructions.

The following example illustrates a Turing/Volta instruction, together with its control information, as decoded by `nvdiasm` [11]:

```
S2UR UR4, SR_CLOCKLO ;
```

The diagram illustrates the decoding of a 128-bit instruction into two 64-bit words. The first word is labeled 'Instruction Part 1' and the second word is labeled 'Instruction Part 2' and 'Control Logic'.

```
/* 0x000000000000479c3 */
/* 0x0000240000005000 */
```

The output shows the instruction decoded into two 64-bit words. The first word contains pure instruction information, while the second contains both instruction and control information.

Our experiments based on instruction disassembly and arbitrary code execution suggest that the encoded 128 bits are used as follows:

- at least 91 bits for instruction information;

¹By control information in this context, we mean instruction scheduling decisions taken by the compiler, that the architecture must enforce; the next section discusses this topic in detail.

- at least 23 bits for control information;
- the remaining 14 bits appeared to be unused in both Turing and Volta; in our experiments, they were ignored both by `cuobjdump` and by the hardware.

2.1 Control information

Control words appeared first with the Kepler architecture, which substantially replaced dynamic hardware scheduling with static software scheduling. Control words encode instruction scheduling decisions taken by the compiler [8] that the hardware must enforce. The design choice to use software scheduling in Kepler was a departure from the previous design (Fermi): designers replaced a complex hardware scheduler with a simpler, more efficient one that occupied less silicon area and consumed less power. Overall, software scheduling enabled simpler on-chip control logic, leading to higher compute density per area of silicon and better energy efficiency.

On Turing and Volta, 128 bits contain one instruction together with the control information associated with only that instruction.

Pre-Volta architectures pack one control word with multiple instruction words into a *bundle*. In each bundle, the first word contains control information, and the remaining words (3 on Pascal and Maxwell, 7 on Kepler) encode one instruction each. Each control word affects how the architecture schedules the instructions within the bundle.

The following excerpt shows a bundle of Pascal instructions decoded by `nvdiasm`. The bundle contains four 64-bit words. The first word, which has a hexadecimal dump but no corresponding disassembled instruction, is a control word. The remaining three words are instructions.

```
/*0288*/      @P5 LDG.E.CI R66, [R86+0x100];      /* 0x000f8800fe2007f1 */
/*0290*/      @!P5 MOV R66, RZ;                  /* 0x0eed4a00010055642 */
/*0298*/      @P6 LDG.E.CI R67, [R86+0x180];      /* 0x0eed4a00018065643 */
```

Control information is encoded as follows on the different GPU generations:

- on Kepler, each control word contains 6 zeroes as its most significant bits, 2 zeroes as its least significant bits, and 7 sections of 8 bits each;
- on Pascal and Maxwell, each control word contains one zero as its most significant bit, and 3 sections of 21 bits each;
- on Turing and Volta, each control section contains 2 zeroes as its most significant bits, and 1 section of 21 bits. For every 128 bits corresponding to one instruction, control information is preceded and followed by bits encoding the instruction itself.

Sections containing control information are organized in the same way on Turing, Volta, Pascal and Maxwell. Each section contains 6 fields, organized as follows:

Width (bits)	4	6	3	3	1	4
Meaning	Reuse flags	Wait barrier mask	Read barrier index	Write barrier index	Yield flag	Stall cycles

Fields have the following meaning:

Reuse flags. Each hardware thread on Turing, Volta, Pascal and Maxwell has a 2-way associative Content-Addressable Memory (CAM) for each of the four conceptual source registers operand positions. This memory is intended to allow data reuse between instructions without accessing any register ports: this relieves pressure on the register file, and helps reducing register bank conflicts (we discuss register bank conflicts at length in Section 3.5.1). Reuse flags control this mechanism as follows: an instruction may flag for saving into the reuse set any combination of up to its first four arguments. Each instruction will attempt to service register reads for its first 4 arguments from the respective reuse slots before resorting to loading values via register file ports. E.g., if the last two reuse-saved registers in the second instruction source operand position were R98 and R99, either of those registers may be used in the second position of instructions without contributing to register bank conflicts. The four bits in the reuse flags map the first to fourth source operands with the least to most significant bits, respectively.

Wait barrier mask; Read/Write barrier index. While most instructions have fixed latency and can be statically scheduled by the assembler, instructions involving memory and shared resources typically have variable latency. Turing, Volta, Pascal and Maxwell use *dependency barriers* to track the completion of variable-latency instructions and resolve data hazards. When a variable-latency instruction writes to a register, the assembler associates it to one of the 6 available barriers by setting the corresponding *write barrier number* field. When a later instruction consumes that register, the assembler marks the instruction as waiting on that barrier by setting the bit corresponding to that barrier in the *wait barrier mask*. The hardware will stall the later instruction until the results of the earlier one are available. An instruction may wait on multiple barriers, which explains why the *wait barrier mask* is a bitmask, not an index.

Read dependency barriers. Read dependency barriers serve to protect against write-after-read hazards. Unbuffered instructions that write the contents of registers to memory need the registers to remain unchanged during the operation. To guarantee that, the assembler associates them to a barrier

by populating the corresponding *read barrier number* field. Later instructions writing to the same register will wait on that barrier.

Stall cycles. This 4-bit field indicates how long the scheduler should wait before issuing the next instruction, ranging from 0 to 15 cycles. On Pascal and Maxwell, if the combination of this field and the yield flag contain a special combination of bits, the two dispatchers in a processing block can dispatch two consecutive instructions of a warp at the same time (dual issue). On Turing and Volta there is only one dispatcher in a processing block, and we do not observe dual issue in the generated code.

Yield flag. As its predecessors, the Turing architecture uses a one-bit yield flag to balance the workload assigned to a processing block. When this bit is set, the scheduler prefers to issue the next instruction from the current warp. When the bit is cleared, the scheduler prefers to switch to another warp, making all register reuse flags for the next instruction ineffective. This costs one extra cycle to switch to another warp.

2.2 Processing Blocks and Schedulers

The Turing streaming multiprocessor (SM) is partitioned into four *processing blocks*, each containing a dedicated warp scheduler and dispatch unit [3]. Instructions from the same warp are allocated to a specific processing block, and

Table 2.1: This experiment reveals the same mapping between warps and schedulers on Turing and Volta: warps with the same index modulo 4 are mapped to the same scheduler. We vary the indices of two active warps (A and B) and measure their aggregate throughput. When the indices collide modulo 4 (i.e., they are mapped to the same scheduler) performance drops. All values are in single-precision GFLOPS.

T4 GPU					
		Warp A Index			
		0	1	2	3
Warp B index	4	48.9	72.4	72.5	73.1
	5	73.4	46.7	72.5	73.1
	6	73.2	72.8	47.0	73.2
	7	72.9	72.7	72.7	46.2
V100 GPU					
		Warp A Index			
		0	1	2	3
Warp B index	4	42.27	66.05	66.04	65.29
	5	66.05	41.98	66.04	66.04
	6	66.02	66.04	42.06	66.04
	7	66.04	66.04	66.02	42.08

can only access the processing units within that block.

We found that warps are mapped to schedulers (and processing blocks) on Turing and Volta according to the same, simple rule:

$$scheduler_id = warp_id \% 4.$$

This is demonstrated with a benchmark composed of 8 warps running on a single SM simultaneously, of which only 2 are active with loops of FFMA instructions, while the remaining 6 are idle.

We repeat the experiments varying the warp index of each of the two active warps (Warp A and B), while measuring each time the aggregate arithmetic throughput achieved by the two warps. The results (see Table 2.1) show that whenever the two warps have the same index modulo 4 (e.g., 0 and 4, 1 and 5, ...), their aggregate performance drops, which suggests that they are mapped to the same scheduler.

These findings are consistent between Turing and Volta.

Furthermore, these results indicate that every block of your workload must use at least 128 threads to fully utilize the processing units on one SM of Turing and Volta.

2.3 Instruction word format

2.3.1 Opcodes

Turing and Volta use more bits to encode their instructions than in previous architectures.

Unlike previous architectures (Pascal, Maxwell and Kepler), which organize the opcode in the most significant bits of the instruction, Turing and Volta architectures place the opcode in the least significant bits of the first 64-bit word of the instruction. Turing opcodes vary in length from 10 to 13 bits. For an extensive opcode reference that compares Pascal, Volta and Turing, see the Appendix.

2.3.2 Operands

As in previous architectures, instruction operands on Turing can be registers of different types, memory addresses (constant, shared or global), or an immediate value. Predication is regulated by 4 bits: the first bit is a negation flag, and the remaining 3 bits encode a predicate register index.

Chapter 3

Memory hierarchy

NVidia GPU architectures tend to increase in complexity with newer generations. Gaining a deep understanding of GPU memory hierarchy as they evolve is necessary to write efficient code.

For designers to map their working sets optimally onto the memory hierarchy, it is especially important to know the size of each cache level, whether that memory is co-located with another cache that might evict its contents,

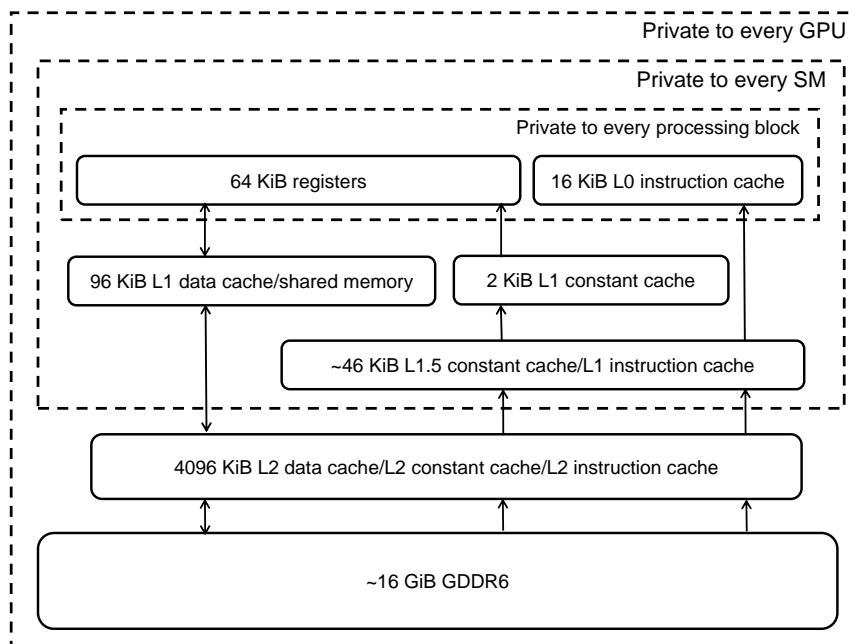


Figure 3.1: Memory hierarchy of the Turing T4 GPU (TU104).

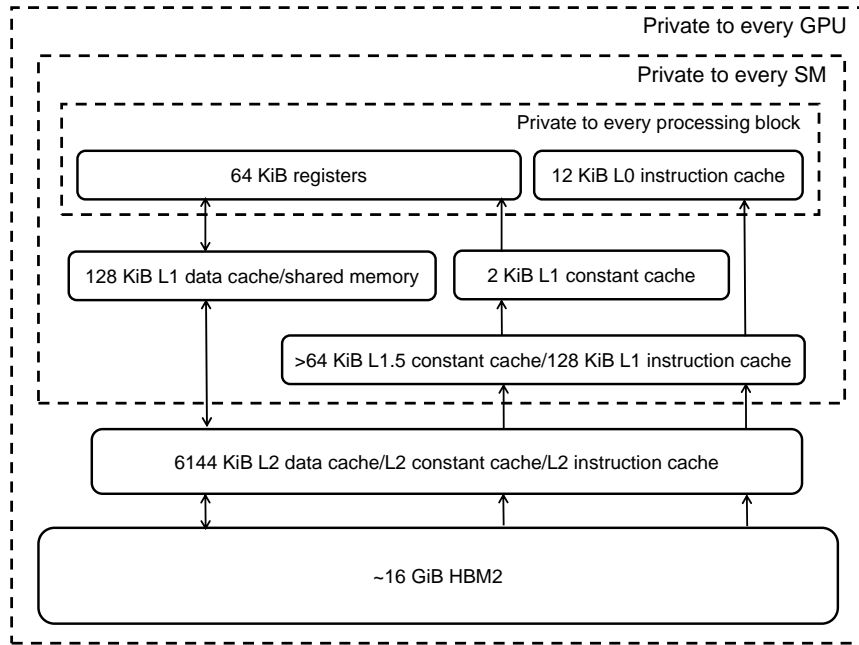


Figure 3.2: Memory hierarchy of the Volta V100 GPU (GV100).

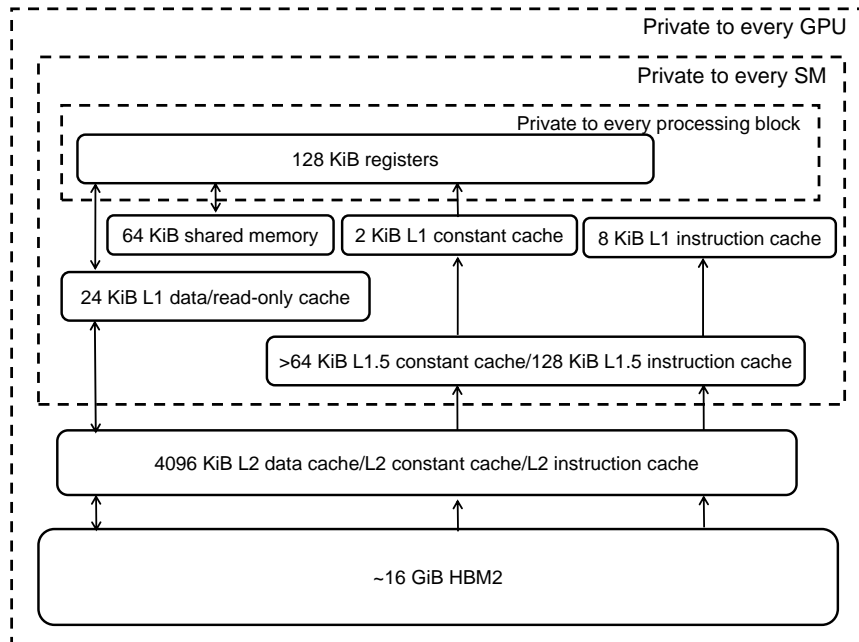


Figure 3.3: Memory hierarchy of the Pascal P100 GPU (GP104).

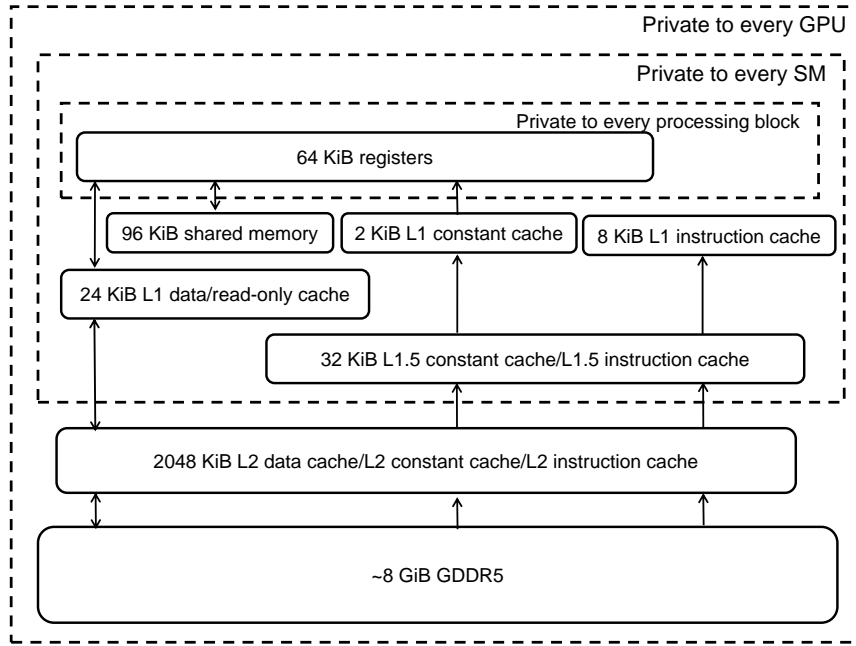


Figure 3.4: Memory hierarchy of the Maxwell M60 GPU (GM204).

and whether each cache memory is private to a streaming multiprocessor or shared among all.

In this chapter, we describe the structure of Turing’s memory hierarchy in detail (Figure 3.1). Specifically, we reveal:

- the geometry, properties and performance of all cache levels and Translation Look-aside Buffers (TLBs);
- register file banks and their conflicts;
- the performance of shared and global memory under load.

Table 3.1 summarizes our findings, also comparing Turing against the Volta, Pascal, Maxwell and Kepler generations.

The T4 GPU employs GDDR6 memory, which offers a bandwidth of 320 GB/s (at a memory clock frequency 5,001 MHz), in conjunction with a L2 cache of 4,096 KiB [3]. Data loaded from global memory is implicitly cached in L1 and L2.

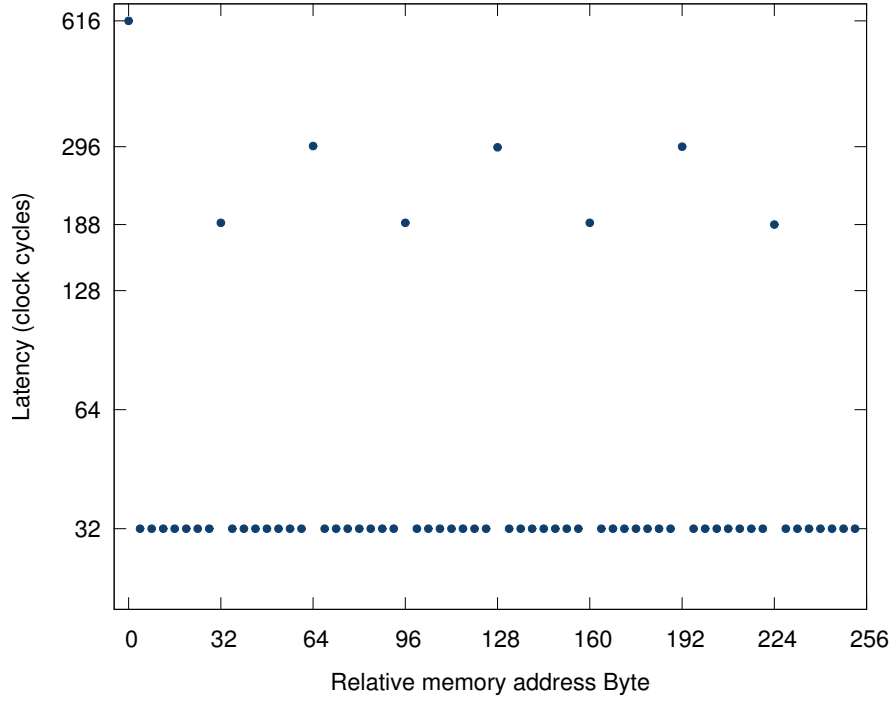


Figure 3.5: Global memory access latency, as per our measurements with the fine-grained p-chase method by Mei and Chu [9]. The 616-cycle latency of the first access is the result of both L2 cache miss and TLB miss. Accesses to the following data, which are stored in the same L1 cache line, enjoy a low L1 hit latency of 32 cycles. Data points with a 188-cycle latency correspond to an L1 miss and L2 hit; points with a 296-cycle latency correspond to an L2 miss and a TLB hit.

3.1 L1 data cache

Turing adopts the same combined L1 data cache / shared memory design as Volta. This design reduces the cache hit latency and improves the bandwidth with respect to the Pascal architecture.

As the geometry of the L1 data cache is concerned, our findings agree with what reported in the Turing and Volta architecture whitepapers [10, 3]. Specifically, the T4 offers twice as much L1 data capacity, and twice as high a bandwidth as the P4 GPU.

As performance is concerned, our experiments show that on a T4 GPU, the L1 data cache offers approximately $3.7\times$ more bandwidth than its P4 predecessor.

3.1.1 Latency and bandwidth

The L1 data cache hit latency we measured on the T4 GPU is 32 cycles, compared to 82 cycles on the P4 (see Figure 3.5).

Before Turing and Volta, Kepler was the most recent architecture to combine its L1 cache and its shared memory. Kepler's L1 cache read hit latency is 35 clock cycles. Turing exhibits a better L1 latency than Kepler in clock cycles, despite the T4 being clocked almost twice as high as the K80 (1,590 vs. 875 MHz).

We use the following benchmark to measure the L1 data cache load bandwidth. The benchmark scans an array with 32-bit elements; every warp accesses all the elements in the array:

```
__global__ void l1_bw( uint32_t *startClk, uint32_t *stopClk,
                     float *dsink, uint32_t *posArray )
{
    // Thread index
    uint32_t tid = threadIdx.x;

    // Side-effect variable, intended to avoid compiler elimination of this code
    float sink = 0;

    // Warm up the L1 cache by populating it
    for (uint32_t i = tid; i < L1_SIZE; i += THREADS_NUM) {
        float *ptr = posArray + i;
        asm volatile ( "{\t\n"
            ".reg .f32 data;\n\t"
            "ld.global.ca.f32 data, [%1];\n\t"
            "add.f32 %0, data, %0;\n\t"
            "}" : "+f"(sink) : "l"(ptr) : "memory"
        );
    }

    // Synchronize all threads
    asm volatile ( "bar.sync 0;" );

    // Start timing
    uint32_t start = 0;
    asm volatile ( "mov.u32 %0, %%clock;" : "=r"(start) :: "memory" );

    // Load data from L1 cache, accumulate
    for (uint32_t i = 0; i < L1_SIZE; i += THREADS_NUM) {
        float *ptr = posArray + i;
        // every warp loads all data in l1 cache
        for (uint32_t j = 0; j < THREADS_NUM; j += WARP_SIZE) {
            uint32_t offset = (tid + j) % THREADS_NUM;
            asm volatile ( "{\t\n"
                ".reg .f32 data;\n\t"
                "ld.global.ca.f32 data, [%1];\n\t"
                "add.f64 %0, data, %0;\n\t"
                "}" : "+f"(sink) : "l"(ptr + offset) : "memory"
            );
        }
    }

    // Synchronize all threads
    asm volatile ( "bar.sync 0;" );

    // Stop timing
    uint32_t stop = 0;
    asm volatile ( "mov.u32 %0, %%clock;" : "=r"(stop) :: "memory" );
}
```

```
// Write time and data back to memory
startClk[tid] = start;
stopClk[tid]  = stop;
dsink[tid]    = sink;
}
```

We report L1 data bandwidths we measured across GPU devices in Table 3.2, together with their theoretical upper bounds.

The actual bandwidth we measure on the T4 GPU is 58.83 bytes/cycle/SM, i.e., $3.7\times$ higher than that of the P4 GPU, i.e., 15.7 bytes/cycle/SM. This bandwidth comparison expressed in cycle counts is meaningful, because the T4 and P4 cards run at very similar graphics clock frequencies f_g .

We calculate the theoretical throughput by multiplying the LSU count per SM by the number of bytes that each LSU can load per cycle per instruction.

Historically, architectures that employ an L1 cache combined with shared memory (Turing, Volta and Kepler) exhibit a higher L1 bandwidth than architectures where the L1 cached and the shared memory are separate (Pascal and Maxwell).

3.1.2 Geometry

According to the Turing whitepaper [3], load/store operations can use a L1 data cache of 32 KiB or 64 KiB in size.

Our experiments based on Mei and Chu’s fine-grained pointer-chase technique [9] were unable to detect the whole configured size, and fell 7 KiB short of the nominal L1 data cache size, on both Volta and Turing architectures (see Table 3.3).

In our experimental setup, the shared memory is configured to a size of 64 KiB. We then employed a benchmark that scans a variable length array **A** twice. As long as the size of **A** exceeds 25 KiB, we detected cache misses.

At this time we are unable to explain this 7-KiB discrepancy. We conjecture it is the result of a newly applied replacement policy that we discuss below. We confirm that it is not associated to the ECC feature (error correction).

Table 3.1 describes the remainder of L1 data cache geometry as we discover it. The line size, load and update granularity of Turing’s L1 data cache are the same as on the Volta, Pascal and Maxwell GPUs.

In our previous report for Volta [2], we discovered an improved L1 cache replacement policy on Volta with respect to its predecessors. Turing also features a L1 cache replacement policy that aims at preserving large arrays from eviction caused by sparse memory accesses.

We employed a benchmark that scans a variable length array twice, and

recorded positions and latency data when L1 cache miss happens. We found that when the L1 data cache saturates, Turing randomly evicts 4 consecutive cache lines (128 B). We observed that once a block of cache lines are evicted, the second scan will cause more cache lines from the same set to be evicted.

3.2 Unified L2 cache

Turing employs an L2 cache that is unified for data, instructions and constant memory, as the previous GPU generations do. The L2 cache on a T4 GPU is a 16-way, set-associative cache having a size of 4,096 KiB, a cache line of 64 B, and an average latency of 188 clock cycles (Figure 3.5).

We use the following benchmark to measure L2 load bandwidth, on all the GPUs considered:

```
__global__ void l2_bw(float *dsink, uint32_t *posArray)
{
    // block and thread index
    UINT tid = threadIdx.x;
    UINT bid = blockIdx.x;

    // accumulator; side effect to prevent code elimination
    float sink = 0;

    // load data from l2 cache and accumulate
    for (UINT i = 0; i < L2_SIZE; i += THREADS_NUM) {
        DTYPE* ptr = posArray + i;
        // every warp loads all data in l2 cache
        for (UINT j = 0; j < THREADS; j += 32) {
            UINT offset = (tid + j) % THREADS;
            asm volatile("{\t\n"
                ".reg .f32 data;\n\t"
                "ld.global.cg.f32 data, [%1];\n\t"
                "add.f32 %0, data, %0;\n\t"
                "}\t\n" : "+f"(sink) : "l"(ptr + offset) : "memory"
            );
        }
    }

    // side effect: store the result
    dsink[tid] = sink;
}
```

Note that we warm up the L2 cache before launching this kernel (code not shown for brevity). The benchmark contains a simple floating-point accumulation into variable `sink`, which is later written to global memory; this accumulation intentionally creates a side effect intended to prevent the compiler from eliminating the entire benchmark code. The marginal cost of this accumulation is negligible with respect to the data access latency.

The bandwidth we measure on the T4 device (see results in Table 3.4) is 30% higher than the P4's, and 59% of the one measured on the larger V100 GPU.

Table 3.1: Geometry, properties and latency of the memory hierarchy across GPU architectures. For consistency, all performance data in this table were measured on PCI-E cards.

Architecture generation			Turing	Volta	Pascal	Pascal	Maxwell	Kepler
GPU Board			T4	V100	P100	P4	M60	K80
GPU Chip			TU104	GV100	GP100	GP104	GM204	GK210
Processors per chip (P)			40	80	56	40	16	13
Max graphics clock (f_g) MHz			1,590	1,380	1,328	1,531	1,177	875
Threads per Multiprocessor			1,024	2,048	2,048	2,048	2,048	2,048
Registers	Number of banks		2	2	4	4	4	4
	Bank width	bits	64	64	32	32	32	32
L1 data	Size	KiB	32 or 64	32...128	24	24	24	16...48
	Line size	B	32	32	32	32	32	128
	Hit latency	cycles	32	28	82	82	82	35
	Load granularity	B	32	32	32	32	32	128
	Update granularity	B	128	128	128	128	128	128
L2 data	Size	KiB	4,096	6,144	4,096	2,048	2,048	1,536
	Line size	B	64	64	32	32	32	32
	Hit latency	cycles	~188	~193	~234	~216	~207	~200
L1 const	Broadcast latency	cycles	~26	~27	~24	~25	~25	~30
	Cache size	KiB	2	2	2	2	2	2
	Line size	B	64	64	64	64	64	64
	Number of sets		8	8	8	8	8	8
	Associativity		4	4	4	4	4	4
L1.5 const	Broadcast latency	cycles	92	~89	~96	~87	~81	~92
	Cache size	KiB	~46	>=64	>=64	32	32	32
	Line size	B	256	256	256	256	256	256
L2 constant	Broadcast latency	cycles	~215	~245	~236	~225	~221	~220
L0 instruction	Cache size	KiB	~16	~12	-	-	-	-
L1 instruction	Cache size	KiB	~46	128	8	8	8	8
L1.5 instruction	Cache size	KiB	-	-	128	32	32	32
L2 instruction	Cache size	KiB	4,096	6,144	4,096	2,048	2,048	1,536
L1 TLB	Coverage	MiB	32	32	~32	~32	~2	~2
	Page entry	KiB	2,048	2,048	2,048	2,048	128	128
L2 TLB	Coverage	MiB	~8,192	~8,192	~2,048	~2,048	~128	~128
	Page entry	MiB	32	32	32	32	2	2
L3 TLB	Coverage	GiB	-	-	-	-	~2	~2
	Page entry	MiB	-	-	-	-	2	2
Shared	Size per SM	KiB	32 or 64	0...96	64	64	96	48
	Size per chip	KiB	1,280 or 2,560	0...7,689	3,584	1,280	1,536	624
	Banks per processor (B_s)		32	32	32	32	32	32
	Bank width (w_s)	B	4	4	4	4	4	8
	LSU count per SM (n_{LSU})		16	32	16	16	32	32
	No-conflict latency	cycles	19	19	24	23	23	26
	Theoretical bandwidth	GiB/s	4,070	13,800	9,519	3,919	2,410	2,912
	Actual bandwidth	GiB/s	3,662	12,080	7,763	3,555	2,122	2,540
	Actual/Theoretical ratio	%	90.9%	87.5%	81.6%	90.7%	88.0%	87.2%
Global	Memory bus		GDDR6	HBM2	HBM2	GDDR5	GDDR5	GDDR5
	Size	MiB	15,079	16,152	16,276	8,115	8,155	12,237
	Max clock rate (f_m)	MHz	5,001	877	715	3,003	2,505	2,505
	Theoretical bandwidth	GiB/s	320	900	732	192	160	240
	Actual bandwidth	GiB/s	220	750	510	162	127	191
	Actual/Theoretical ratio	%	68.8%	83.3%	69.6%	84.4%	79.3%	77.5%

Table 3.2: L1 cache load throughput per SM.

	T4	V100	P100	P4	M60	K80	
Theoretical upper bound	64.0	128.0	64.0	64.0	128.0	128.0	bytes/cycle
Measured throughput	58.8	108.3	31.3	15.7	15.7	68.6	bytes/cycle

Table 3.3: Detectable L1 data cache size with the pointer-chase benchmark on the T4 GPU.

Configured size of shared memory (KiB)	32	64
Expected size of L1 data cache (KiB)	64	32
Detected size of L1 data cache (KiB)	57	25

Table 3.4: L2 data cache load throughput.

	Turing T4	Volta V100	Pascal P100	Pascal P4	Maxwell M60	Kepler K80
Throughput (GB/s)	1,270	2,155	1,624	979	446	339

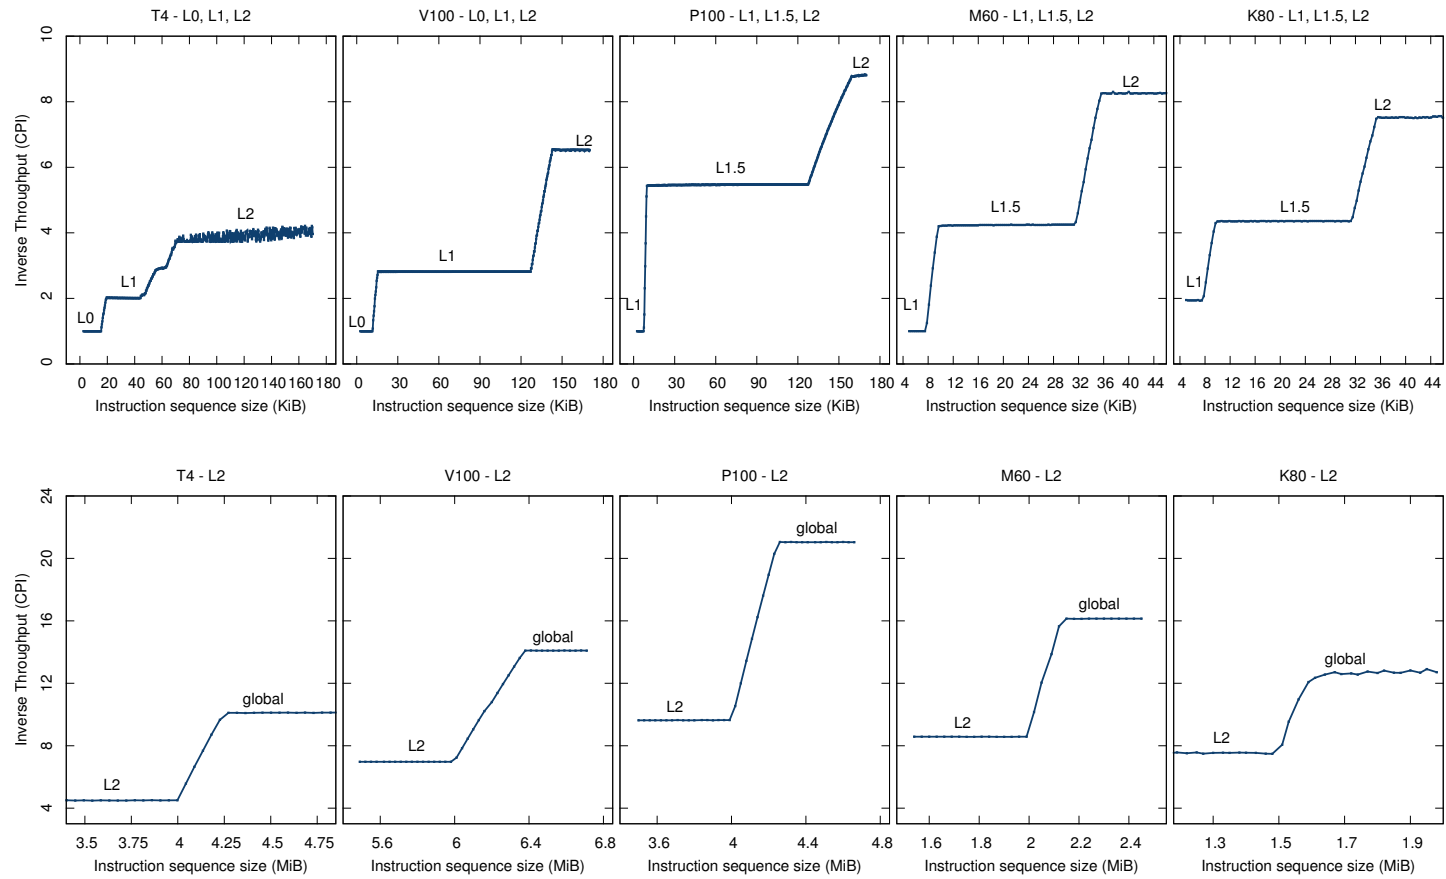


Figure 3.6: We detect the size of instruction cache level with a benchmark based on sequences of identical instructions of increasing length. We then chart the average inverse throughput: each plateau reveals the size of a cache level. **Top charts:** boundaries of the first two levels in the hierarchy. **Bottom charts:** limits of the last level, misses to global memory.

3.3 Instruction cache hierarchy

In this section, we map experimentally the size and the organization of the instruction cache hierarchy. In practice, that consists in (1) detecting the size of each cache level and (2) determining how cache levels are distributed within the architectural blocks (scheduler, SM, entire chip) of the GPU.

3.3.1 Taxonomy

All GPU architectures we considered, including Turing, feature three levels of instruction caches. To avoid confusion, note that on Turing and Volta the three levels are named differently (L0, L1, L2) than on previous architectures (L1, L1.5, L2). We adopt this established taxonomy for consistency with the NVidia’s whitepapers [3, 10] and with prior literature. Pay attention to expressions like “the second level of instruction caches”: this expression refers to L1 on Turing and Volta, but to L1.5 on Pascal, Maxwell and Kepler.

3.3.2 Size

To detect the size of each cache level, we study how the average inverse throughput (i.e., average clocks per instruction, or CPI) achieved by a long sequence of instructions changes as a function of sequence length. As we increase the length of a sequence, we expect to see a constant CPI value until the sequence exceeds the cache size. Indeed, experimental results show plateaus and ramps (Figure 3.6) which correspond to cache level sizes and transitions from one level to the following. In the figure, the bottom charts focus on the three instruction cache levels, whereas the bottom charts focus on the transition between the last cache level and global device memory.

We report all findings in Table 3.1. Turing enjoys better inverse throughput than its predecessors when accessing the second and third instruction cache levels.

Experimental setup. Our benchmark measures the average CPI seen by a sequence of instructions of given length that exert no pressure on the data cache hierarchy. We iterate measurements for sequence sizes starting from the cache line size up to the larger plausible size of L3. The benchmark executes each sequence twice, but only times the second execution, so that we only measure *capacity misses* and *conflict misses*, but not *cold misses*.

- On Pascal, Maxwell and Kepler, we employ the same technique as in our previous report [2] for the sake of consistency, i.e., long sequences of FFMA instructions, whose register operands are chosen so that each instruction experiences no register dependence with its neighbors.
- On Volta and Turing, we switched to a simpler method that uses NOP sequences rather than FFMA. This choice circumvents NVCC’s undesired generation of 2-cycle stalls between consequent FFMA instructions on these two GPUs.

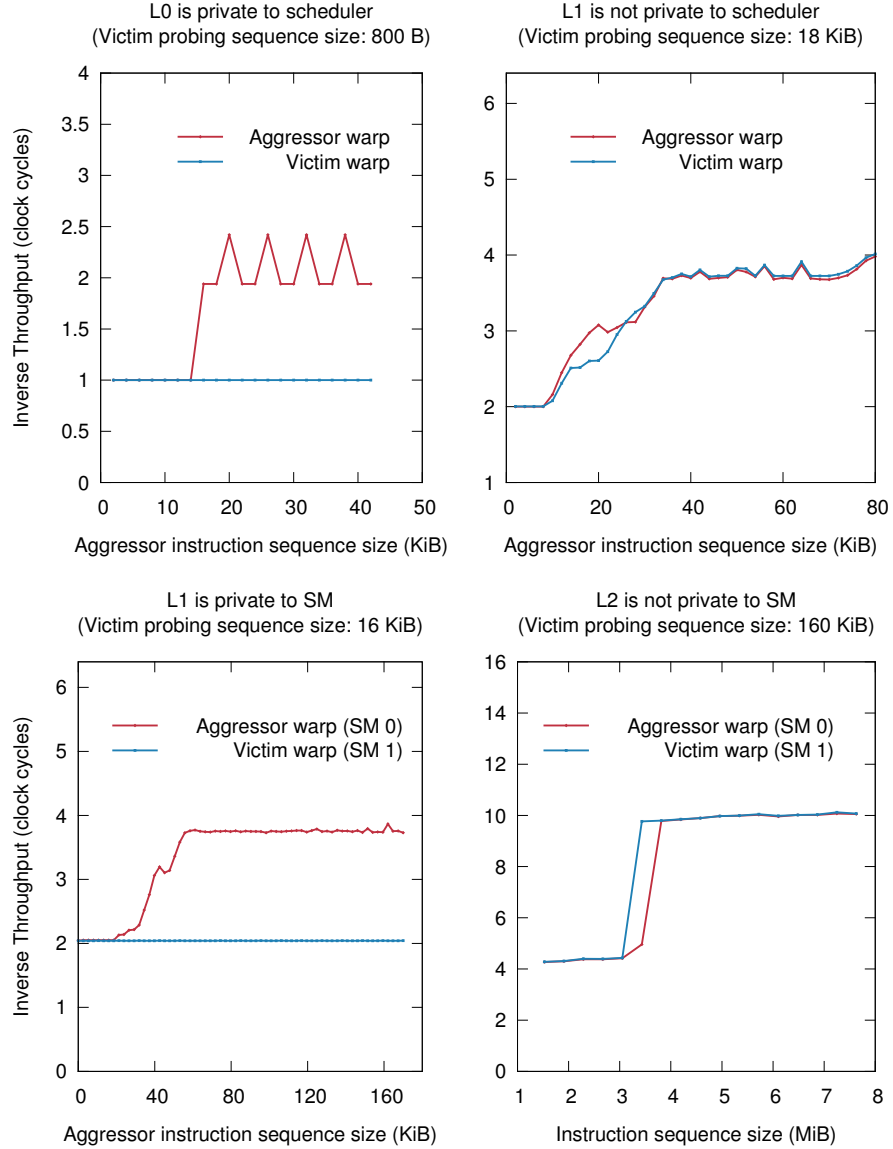


Figure 3.7: Aggressor-victim experiments designed to detect what architectural block (a scheduler, an SM, the entire chip) owns each level of the instruction cache, by observing how an aggressor warps that causes instruction cache pressure affects a victim warp's performance. **Top left:** each L0 instruction cache is private to a scheduler. **Top right:** an L1 instruction cache is not private to a scheduler. **Bottom left:** each L1 instruction cache is private to one SM. **Bottom right:** the L2 instruction cache is common among all SMs.

3.3.3 Organization

Across the different GPU architectures, levels in the instruction memory hierarchy are organized as follows:

- on Turing and Volta, each L0 instruction cache is private to one scheduler/processing block;
- on all GPUs considered, each L1 instruction cache is private to an SM;
- on Pascal, Maxwell and Kepler each L1.5 instruction cache is private to one SM; the L1.5 instruction cache does not exist on Turing and Volta;
- on all GPUs considered, the L2 cache is unified (i.e., it caches instructions and data) and it is shared across all SMs.

On architectures older than Turing, we provided experimental support for these claims in our previous report [2]. For claims about Turing, we collected evidence using experiments designed as follows.

Our experiments measure the interaction between an *aggressor warp* and a *victim warp*. Both warps loop through sequences of NOP instructions of chosen length:

- the **victim** warp only runs a fixed-length NOP sequence, typically designed to fit within a certain instruction cache level; we call it the *victim probing sequence*;
- the **aggressor** warp runs, in addition to the same probing sequence as the victim, and *before* it, a variable-length NOP sequence, designed to thrash a given cache level, potentially evicting instruction cache entries.

We monitor whether the evictions caused by the aggressor warp only affect its own performance, or they affect the victim as well: if the victim is unaffected, then the smallest cache level that fits the fixed-length victim probing sequence is private to the architectural block where the two warps are running (i.e., GPU processing block or SM); else, the cache level is shared between the two warps and located outside the block considered. In our experiments, both warps monitor their performance by measuring their inverse throughput (CPI).

Results show that each L0 instruction cache is private to a processing block, that each L1 instruction cache is private to an SM, and that the L2 cache is shared among all SMs (Figure 3.7).

To examine the relation between levels L0, L1 and **schedulers** (or GPU processing blocks), we use experiments where the aggressor and victim warps run on the same SM, but different processing blocks. We use increasingly longer sequences in the aggressor warp. To exclude compulsory misses from

the measurements, we let the aggressor and then the victim warm up the caches by running each their respective sequence once.

We observe that:

- as the aggressor sequence grows while remaining below L0 capacity, only the aggressor experiences a slowdown (top left chart in Fig. 3.7), whereas the victim is unaffected. This indicates that the two warps access distinct L0 caches, private to each processing block;
- as the instruction sequence grows above L0 capacity (top right chart) and into L1, both warps slow down similarly, which indicates that the two warps share L1.

Next, we examine the relation between levels L1 and L2, and SMs, with similarly constructed experiments. This time, the two warp run on separate SMs (SM0 and SM1).

We observe that:

- as the aggressor sequence exceeds L0 but remains within L1 capacity, only the aggressor warp experiences a slow-down corresponding to L1 hit rates (bottom left); the victim, still running a sequence fitting L0 (16 KiB), is unaffected. This indicates that different SMs have distinct L1 caches;
- as the aggressor sequence exceeds L2 capacity (bottom right chart), both victim and aggressor experience slowdowns; This indicates that different SMs access the same L2 cache.

3.4 Constant memory hierarchy

The constant memory is a cached window of global memory, reserved for data declared with the `__constant__` keyword, plus kernel invocation parameters and immediate constants. We find that Turing has three levels of constant cache memory, which have the geometry and properties described in Table 3.1 and latency as in Figure 3.9.

The constant memory hierarchy used in Turing did not change significantly from previous generations. Across all the GPU generations we considered, the following properties hold true:

- the L1 constant cache uses a non-LRU replacement policy;
- each SM possesses two private levels of constant caches, which we denote as L1 and L1.5 constant cache (accesses to either of each level within an SM do not affect the same cache levels on other SMs);




Figure 3.8: An aggressor-victim experiment shows that the L1.5 constant cache and the L1 instruction cache coincide. We measure the miss rates experienced by the scan of a constant array pre-cached in constant L1.5 cache (victim) that follows a long sequence of identical `FFMA` instructions (aggressor), intentionally designed to cause L1 instruction cache pressure. As the aggressor’s sequence length increases, the victim suffers increasing miss rates.

- the L2 cache is the third level of constant cache. It is shared among all SMs and is unified for instruction and data.

On Turing as in Volta, the second levels of the constant and the instruction cache are backed by the same hardware cache. More precisely, the L1.5 constant cache and the L1 instruction cache coincide. To prove this claim, we run an aggressor-victim experiment, in which we show that instruction sequences of increasing length (aggressor) evict pre-populated entries in the L1.5 constant cache. We detect these evictions by recording the execution time of a constant array scan (victim) that we execute after the aggressor. We use instruction sequences composed of identical `FFMA` instructions.

Experimental results (Figure 3.8) show that longer instruction sequences in the aggressor cause correspondingly higher miss rates in the victim. We observed victim miss rates vary from 0% to 100%.

As in previous architectures, constant memory accesses on Turing support




Figure 3.9: Latency of concurrent loads from constant memory within a warp depends on where the data is found in the cache hierarchy (L1, L1.5, or L2) and on the count of distinct locations referenced. The hardware broadcasts accesses to the same location.

broadcasting (see Figure 3.9). When all threads within a warp access the same address, the constant memory sends data to all threads simultaneously. When threads visit diverging addresses, the accesses are serialized.

3.5 Registers

3.5.1 Register File Banks

Turing and Volta use a physical register file of 16,384, 32-bit elements in each processing block. Thread-visible logical registers are allocated in increments of 8, or aggregate increments of 256 for all 32 threads in a warp. These register files are organized in 2 banks with dual 32-bit ports each, with logical registers belonging to the bank with the index matching their name, modulo-2. Each port can satisfy only one 32-bit read per clock cycle, and instructions in compiled code requiring three or more operands (such as `FFMA`, the single-precision floating-point fused multiply-and-add instruction) will suffer a stall in execution due to a *register bank conflict* if any three source registers' names map to either dual-ported bank.




Figure 3.10: Register bank conflicts affect the execution time of instructions. Charted is the execution time taken by long sequences of identical `FFMA` instructions, as we vary one source register (`RX`). In both sequences `R6`, the destination sequence, is irrelevant. In sequence `FFMA R6, R97, R99, RX`, the choice of `RX` causes a conflict when `RX` is odd: the other two source operands are already using both ports from bank 1, and a third access cannot occur in the same clock cycle. In sequence `FFMA R6, R98, R99, RX`, no choice of `RX` can cause a conflict because `R98` and `R99` are on different banks.

For example:

- instruction `FFMA R15, R11, R12, R13` has no conflict, since source operands `R11` and `R13` can be serviced by bank 1's two ports, `R12` can be serviced by one of bank 0's ports, and destination register `R15` does not use an additional port from bank 1;
- instruction `FFMA R18, R10, R12, R16` suffers a conflict because `R10`, `R12` and `R16` are all in bank 0. (The destination of `R18` is irrelevant.)

Architectures prior to Volta used 4, single-ported banks, requiring substantially more constrained register scheduling by the compiler, but there are opportunities for improvements even on the newest devices. In our technical report on Volta [2], we demonstrated performance increases of up to 15% by minimizing bank conflicts through careful register re-assignment.

Figure 3.10 illustrates the effect of register bank conflicts on instruction latency on the T4 GPU. We use long sequences of identical `FFMA` instructions in which we vary one source register index (`RX`) to cause conflicts. Since the T4 GPU has dual-ported register banks, a conflict will only happen when all three 32-bit source registers in an `FFMA` instruction belong to the same bank. In every instruction of form `FFMA R6, R97, R99, RX` in the benchmark, `R97` and `R99` are in bank 1; if `RX` also sits in bank 1, a conflict will occur. (`R6` is irrelevant as it is a destination register.) In instruction sequence `FFMA R6, R98, R99, RX`, because `R98` and `R99` sit in different banks, there is no choice of `RX` that can cause three reads from the same bank.

3.5.2 Uniform Registers

As per NVIDIA’s documentation, Turing introduces a new feature intended to improve the maximum achievable arithmetic throughput of the main, floating-point capable datapaths, by adding a separate, integer-only, scalar datapath (named the *uniform datapath*) that operates in parallel with the main datapath.

This design is intended to accelerate numerical, array-based, compute-bound workloads that occupy the main datapaths almost completely with floating-point instructions, typically `FFMA` or `HMMA`, but also contain a few integer operations, typically updating array indices, loop indices or pointers; or performing array or loop boundary checks.

These few integer instructions spoil the instruction mix, and prevent the main datapaths from ingesting a 100% pure stream of `FFMA` or `HMMA`. In these circumstances, even a small fraction of integer instructions can hurt the overall arithmetic throughput, lowering it significantly from its theoretical maximum.

On Turing, the compiler has the option to push these integer operations onto the separate uniform datapath, out of the way of the main datapath. To do so, the compiler must emit *uniform datapath instructions*.

Regular instructions can access both uniform and regular registers. Uniform datapath instructions, instead, focus on uniform instructions almost exclusively.

While at this time we have not been able to stimulate the generation of uniform datapath instructions by the compiler, we were able to enumerate the 64 uniform registers supported by Turing (including a Uniform Zero Register `URZ` and 63 general-purpose uniform registers `UR0–UR62`) by systematically disassembling packed uniform instructions.

3.5.3 Regular Registers

Instructions on Turing still supports the 256 regular registers (including the general-purpose `R0–R254` and the Zero Register `RZ`).




Figure 3.11: Shared memory latency increases under contention. Both axes use exponential scales. We used a stride factor multiplying the thread index as an offset to load data from shared memory. Each thread visits one 32-bit element and measures the average access latency. The benchmark warms shared memory before recording clock cycles.

We found that the `cuobjdump -dump-resource-usage` command (that prints a kernel’s register usage) reports a count that includes both regular and uniform registers. The upper limit of total registers used in any CUDA kernel is 256, unchanged from Volta.

We confirmed this result by patching the register count in the section header of a CUDA kernel to values above 256, and determining that `cuobjdump` only recognizes 256 registers at most.

3.6 Shared memory

The T4 GPU has up to 64 KiB of shared memory (configurable by the user) that offers low latency and high memory bandwidth. In this section, we characterize shared memory performance, including performance under contention.




Figure 3.12: Theoretical and measured shared memory bandwidth on the considered GPUs. The theoretical limits are given by the minimum of product $P \cdot B_s \cdot w_s \cdot f_g$ and product $P \cdot B_s \cdot n_{LSU} \cdot f_g$. The meaning of all factors in these products is explained in Table 3.1.

3.6.1 Latency

Turing’s shared memory enjoys relatively low latency among the GPUs we examined (Figure 3.11). Only the V100 GPU exhibits lower shared memory latency than the T4 GPU.

On all GPUs except for Kepler, the measured average access latency monotonically increases with the number of conflicts in a warp. Kepler is the only GPU adopting dual-ported shared memory banks, allowing any two threads to alias on any given bank without penalty and resolving two further aliases at a time for conflicted banks.

3.6.2 Bandwidth

Due to their large number of streaming multiprocessors, the V100 and P100 GPUs provide the highest theoretical and measured shared memory bandwidth (Figure 3.12).

As benchmarking is concerned, on Kepler, Maxwell, Pascal and Volta, we

were able to rely on `nvprof` to collect shared memory metrics. On Turing, because `nvprof` does not support shared memory metrics collection on that GPU, we resorted to adopting the following custom-tailored benchmark:

```
// Pointer-chasing shared memory bandwidth benchmark
// dData  : Pointer-chase array
// dSink  : Side-effect destination variable (prevents code elimination)
// repeat : Count of pointer-chase steps requested

// To ensure all LSUs in an SM are used, use >= 128 threads
#define THREAD_NUM 1024

// shared memory per block
#define PCHASE_SIZE 8*THREAD_NUM

__global__ void bandwidthTest(uint32_t * dData,
                             uint32_t * dSink,
                             uint32_t  repeat){
    // Pointer-chase starting position in shared memory
    uint32_t sid = threadIdx.x;

    // The pointer-chase array in shared memory
    __shared__ DTYPE shrData[PCCHASE_SIZE];

    // Initialize the pointer-chase array in shared memory
    for (uint32_t i = sid; i<PCCHASE_SIZE; i+=THREAD_NUM)
        shrData[i] = dData[i];

    // Synchronize threads in a same block
    __syncthreads();

    // Scan the shared-memory array with the p-chase method
    unsigned next=sid;
    for (uint32_t j = 0; j < repeat; j++) {
        next = shrData[next];
    }

    // Side effect to prevent the compiler from eliminating this code
    dSink[sid] = next;
}
```

This benchmark performs pointer-chase accesses to the shared memory with a varying number of steps. We invoke as many threads and blocks as possible to provide enough pressure on load/store units. We measured the execution time as we increased pointer-chase step count.

We cross-verified the correctness and accuracy of this benchmark by running it on all architectures other than Turing (on which shared memory metrics are supported) and confirming that the bandwidths it measures match those computed from `nvprof` metrics.

3.7 Global memory

We measured the actual global memory bandwidth and compared it against its theoretical limit for all the GPUs considered (Figure 3.13).

Thanks to their adoption of HBM2 memory, V100 and P100 boards feature a significantly higher bandwidth than GPUs based on GDDR memory. The




Figure 3.13: Theoretical and actual global memory bandwidth on all GPUs considered. Theoretical bounds are derived from NVidia’s whitepapers. Actual bandwidths are the results of our benchmark, which loads data from a global memory array and stores it into another global memory array.

P100 outperforms GDDR-based GPUs boards but suffers from a large gap between actual and theoretical performance. Compared to the P4 GPU, the T4 GPU enjoys a higher global bandwidth because of GDDR6 memory. However, the actual-to-theoretical bandwidth ratio on the T4 board is lower than on the P4 board (68.8% vs. 84.4%)

3.8 TLBs

On Turing and on all other architectures that we examined, we found that

- the L1 data cache is indexed by virtual addresses, and
- the L2 data cache is indexed by physical addresses.

Because L2 is a physical cache, accesses to it involve the TLBs. We prove this claim by scanning a large array with L1 data cache enabled; we size the array to exceed the L1 TLB coverage, so that accesses in the benchmark would




Figure 3.14: Global memory access latency seen by the pointer chase benchmark as it sweeps TLBs. The benchmarks perform a traditional pointer chase after a TLB warm-up scan, calculating the average global memory access latency with a stride of TLB page entry size.

cause at least one level of TLB miss if L1 data cache were indexed by physical address. As expected, we saw no TLB misses in the second scan, as long as the stride is big enough to cache all accesses in L1 data cache. The same benchmark shows that addressing data in L2 data cache goes through the TLBs when the L1 data cache is disabled.

Figure 3.14 shows that, within the available global memory size, there are two levels of TLB on the Turing GPUs. The L1 TLB has 2 MiB page entries and 32 MiB coverage. The coverage of the L2 TLB is about 8192 MiB, which is the same as Volta.

Chapter 4

Instruction latency and throughput

In this chapter, we report on the latency of native Turing instructions. We also benchmark the performance of atomics operations on Turing and compare it with that of older devices. We evaluate the floating-point performance in single, double and half precision on a T4 GPU, and evaluate the updated Tensor Cores.

4.1 Native instructions

Turing and Volta’s instructions typically exhibit lower latency than Pascal and older GPU generations, but Turing does not seem to offer instruction latency improvements over Volta. In this section, we report the latency of common instructions on Turing, Volta and Pascal in Table 4.1.

As the Turing whitepaper [3] mentions, the dependent-issue latency for core FMA math operations is 4 clock cycles, the same as on Volta.

On Turing, we found that most integer, single- and half-precision instructions have similar latencies as those on Volta, whereas double-precision instructions increased their latency above 40 cycles.

On Volta, most integer and single-precision instructions have a latency of 4 cycles. In our previous work we determined that most Volta double-precision instructions have a latency of 8 cycles, and half-precision instructions have a latency of 6 cycles.

On Maxwell and Pascal, instructions IMAD and IMUL have a long latency because they are emulated.

On Pascal, most integer and single-precision instructions have a latency of 6 cycles; double-precision instructions have a latency of 8 cycles; more complex instructions, some of which run on the SFU, require 14 cycles.

Experimental setup. Measuring *dependent issue* instruction latency on a software-

Table 4.1: Latency of frequently used instructions on Volta and Pascal.

Architecture	Instructions	Latency (cycles)
Pascal	BFE, BFI, IADD, IADD32I, FADD, FMUL, FFMA, FMNMX, HADD2, HMUL2, HFMA2, IMNMX, ISCADD, LOP, LOP32I, LOP3, MOV, MOV32I, SEL, SHL, SHR, VADD, VABSDIFF, VMNMX, XMAD	6
	DADD, DMUL, DFMA, DMNMX	8
	FSET, DSET, DSETP, ISETP, FSETP	12
	POPC, FLO, MUFU, F2F, F2I, I2F, I2I	~14
	IMUL, IMAD	~86
Volta	IADD3, SHF, LOP3, SEL, MOV, FADD, FFMA, FMUL, ISETP, FSET, FSETP	4
	IMAD, FMNMX, DSET, DSETP	5
	HADD2, HMUL2, HFMA2	6
	DADD, DMUL, DFMA	8
	POPC	~10
Turing	FLO, BREV, MUFU	~14
	IADD3, SHF, LOP3, SEL, MOV, FADD, FFMA, FMUL, ISETP, FSET, FSETP	4
	IMAD, FMNMX, DSET, DSETP	5
	HADD2, HMUL2, HFMA2	6
	POPC, FLO, BREV, MUFU	~15
	DADD, DMUL	~48
	DFMA, DSET, DSETP	~54

scheduled GPU requires the use of custom-tailored benchmarks designed as follows. To measure the latency of instruction A, we add a second instruction B that depends on A, then set the control word that regulates A's execution:

- if A has fixed latency, we choose a B that consumes A's output. We decrease A's stall cycles in its control word, till A's result consumed by B is incorrect. The last stall value producing correct results is A's latency;
- if A has variable latency, we choose a B of known latency, then set control flags to create an artificial read/write dependency between A and B. We let the scheduler wait for the dependency, then measure the pair's cumulative latency with a bracket of CS2R instructions, and obtain A's latency by subtracting B's known one.

4.2 Atomic operations

Our measurements show that atomic operations on shared memory have a slightly longer latency on Turing than on Volta, but shorter than Pascal and older generations. In Table 4.2, we report those latencies expressed in clock

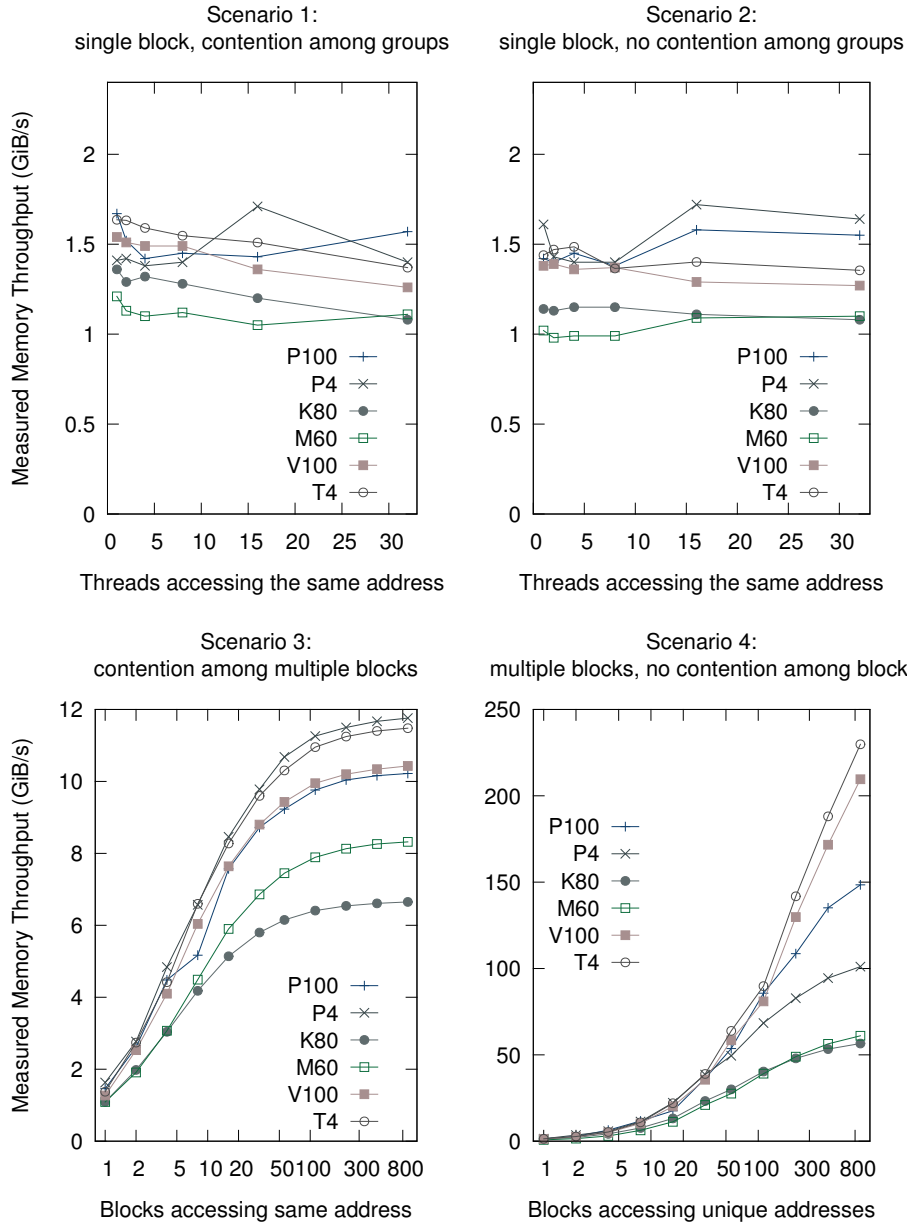


Figure 4.1: Throughput of `atomicAdd` operations on global memory, measured in four contention scenarios.

Table 4.2: Latency of atomic operations on shared and global memory, in clock cycles.

Contention	Shared memory						Global memory					
	T4	V100	P100	P4	M60	K80	T4	V100	P100	P4	M60	K80
none	8	6	15	16	17	93	76	36	26	30	24	29
2 threads	10	7	17	18	19	214	72	31	31	50	26	69
4 threads	14	11	19	25	25	460	73	32	48	50	41	96
8 threads	22	18	30	30	31	952	81	41	48	51	41	152
16 threads	37	24	46	46	47	1,936	97	58	50	51	46	264
32 threads	69	66	78	78	79	4,257	116	76	50	51	46	488

cycles. The comparison is meaningful even in real terms because the different GPUs adopt similar clock frequencies (reported in Table 3.1)

As atomics on global memory are concerned, latency seems to have increased on the T4 device compared with V100. The M60 GPU had the best latency among all GPU considered.

Notably, Kepler is the only architecture where shared memory atomics are slower than global memory one, and by a large margin ($4\times$ to $8\times$). This is due to Kepler’s lack of hardware support for shared memory atomics. Moreover, its emulated atomics degrade quickly under contention. Later architectures support atomics in hardware, and offer low-latency atomics, even in presence of contention.

We measured these latencies with benchmarks designed in the following manner: we determine the latency of atomic instruction A by following it with a load instruction B, of known latency, that visits the same location. We deduce A’s latency from that of pair (A,B) as described in the previous section.

Figure 4.1 reports the throughput measured on GPUs from Kepler to Turing in presence of contention, in four scenarios:

- *Scenario 1*, one block of 1,024 threads. Of these, R threads access the same address, while the others access distinct, sequential addresses in global memory. 8 groups of threads access the same L2 cache line;
- *Scenario 2*, one block of 1,024 threads. Of these, R threads access the same address, while the others access sequential L2 cache lines in global memory, with every group of threads accessing a single L2 cache line;
- *Scenario 3*, a variable number of blocks, of 1,024 threads each. All threads in all blocks access the same address; heavy contention exists among blocks;
- *Scenario 4*, a variable number of blocks, of 1,024 threads each. All threads within a block access the same address. Different blocks access distinct

addresses; no contention exists among blocks.

The T4 GPU doesn't achieve the highest throughput in the scenarios with contention and the scenarios on single SM. The only scenario in which the T4 GPU provides the best performance is on multiple SMs and without contention among SMs. In all scenarios, from Maxwell to Pascal the aggregate throughput increase substantially.

4.3 New Tensor Core instructions

The Turing architecture refreshes its Tensor Cores by offering support for a wider range of operand types than Volta. Specifically, Tensor Cores as introduced in Volta were designed to offer high throughput when performing matrix math on half-precision floating point operands; on Turing, Tensor Cores add support for short integer operands: `int8`, `int4` and `int1`.

Moreover, Turing offers new instructions that allow to express matrix math more succinctly. To demonstrate that, we will compare the Volta and the Turing code generated by the compiler for the same warp-level primitive `wmma::mma_sync()`. Readers will recognize this example from Chapter 4.3 of our technical report on Volta [2].

When targeting Volta, NVCC compiles one example invocation of the primitive into the following 16 `HMMA.884.F32.F32.*` instructions:

```
HMMA.884.F32.F32.STEP0 R8, R26.reuse.COL, R16.reuse.COL, R8 ;
HMMA.884.F32.F32.STEP1 R10, R26.reuse.COL, R16.reuse.COL, R10 ;
HMMA.884.F32.F32.STEP2 R4, R26.reuse.COL, R16.reuse.COL, R4 ;
HMMA.884.F32.F32.STEP3 R6, R26.COL, R16.COL, R6 ;

HMMA.884.F32.F32.STEP0 R8, R20.reuse.COL, R18.reuse.COL, R8 ;
HMMA.884.F32.F32.STEP1 R10, R20.reuse.COL, R18.reuse.COL, R10 ;
HMMA.884.F32.F32.STEP2 R4, R20.reuse.COL, R18.reuse.COL, R4 ;
HMMA.884.F32.F32.STEP3 R6, R20.COL, R18.COL, R6 ;

HMMA.884.F32.F32.STEP0 R8, R22.reuse.COL, R12.reuse.COL, R8 ;
HMMA.884.F32.F32.STEP1 R10, R22.reuse.COL, R12.reuse.COL, R10 ;
HMMA.884.F32.F32.STEP2 R4, R22.reuse.COL, R12.reuse.COL, R4 ;
HMMA.884.F32.F32.STEP3 R6, R22.COL, R12.COL, R6 ;

HMMA.884.F32.F32.STEP0 R8, R2.reuse.COL, R14.reuse.COL, R8 ;
HMMA.884.F32.F32.STEP1 R10, R2.reuse.COL, R14.reuse.COL, R10 ;
HMMA.884.F32.F32.STEP2 R4, R2.reuse.COL, R14.reuse.COL, R4 ;
HMMA.884.F32.F32.STEP3 R6, R2.COL, R14.COL, R6 ;
```

When targeting Turing, NVCC compiles the same primitive invocation into only 4 `HMMA` instructions of a new kind, that contain the new `.1688` infix:

```
# Turing rendition
HMMA.1688.F32 R8, R12, R22, R8 ;
HMMA.1688.F32 R4, R12, R23, R4 ;
HMMA.1688.F32 R8, R2, R24, R8 ;
HMMA.1688.F32 R4, R2, R25, R4 ;
```

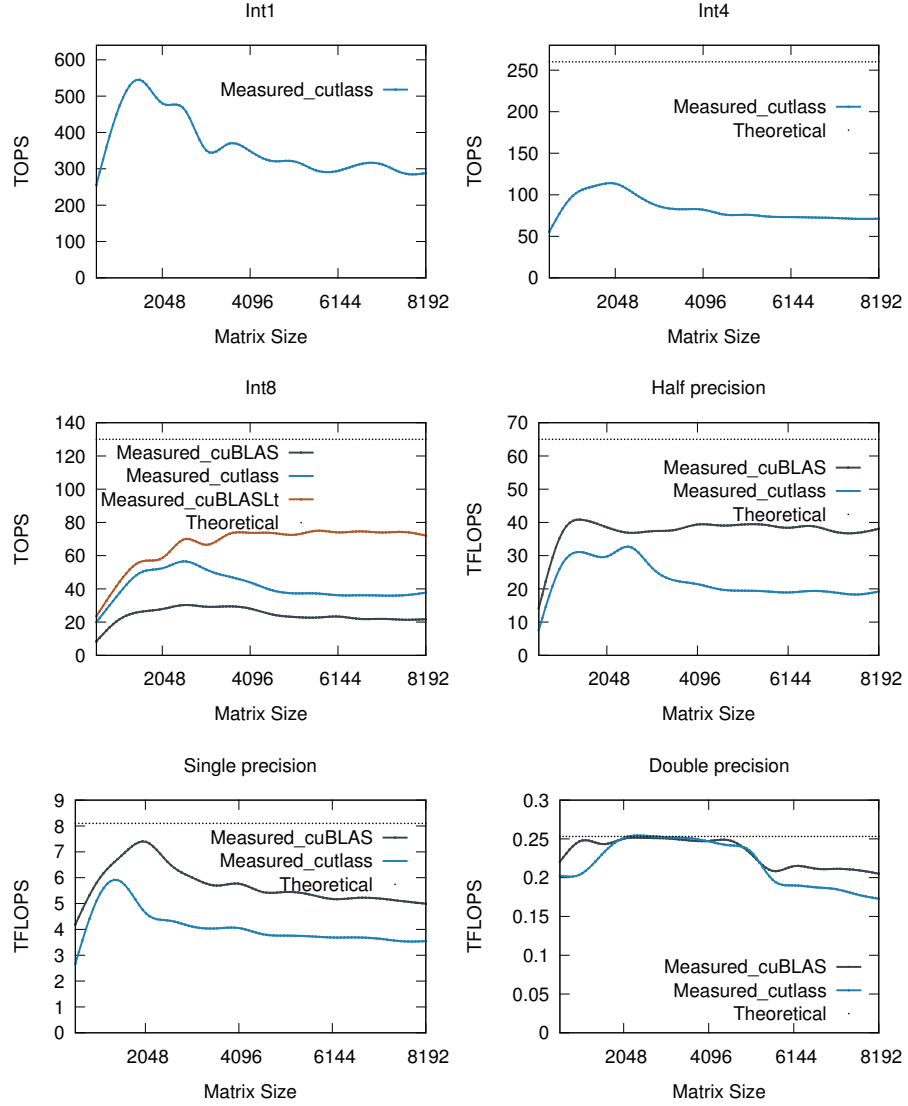


Figure 4.2: Floating-point performance of cuBLAS and CUTLASS matrix multiplication on a T4 GPU running at 1,590 MHz.

4.4 Arithmetic performance

We evaluated arithmetic performance by benchmarking matrix-matrix multiplications using functions from the cuBLAS 10.1 library and template functions from cutlass 1.2, on integer operands and floating-point ones of different precisions. We report arithmetic throughput in TOPS and TFLOPS, when op-

Table 4.3: Arithmetic throughput of matrix multiplication on inference-oriented GPUs on floating point and integer types.

	T4	P4	
Double precision	253	231	GFLOPS
Single precision	7,174	6,944	GFLOPS
Half precision	41,616	6,571	GFLOPS
Int8 precision	74,934	24,172	GOPS
Int4 precision	114,384	-	GOPS
Int1 precision	552,230	-	GOPS

erating on integer and floating-point values respectively. In all experiments, the T4 GPU was running at a clock frequency of 1,590 MHz.

In half, single and double precision, cuBLAS provides higher arithmetic throughput than cutlass. This is because the cuBLAS library has been specifically optimized for the Turing architecture. For `int8` precision, two APIs are available in cuBLAS 10.1:

- BLAS-like extension function `cublasGemmEx`, which invokes native CUDA core implementations, and
- the new light-weight `cublasLtMatmul` function, which supports `int8` native TensorCore implementations.

For `int8`, the throughput of (`cublasLtMatmul`) is much higher than the throughput of (`cublasGemmEx`). At the time of this writing, only cutlass supports `int4` and `int1` matrix multiplication on NVidia GPUs.

Except in double precision, benchmarks don't achieve near-peak performance. For `int8` and `int4`, cutlass implementations don't achieve 50% of theoretical throughput on the T4 GPU (Figure 4.2).

In Table 4.3 we compare the arithmetic throughputs achieved on T4 and P4 GPUs on matrix multiplication at different precisions, with both boards running at the respective top frequencies (1,590 and 1,531 MHz). The T4 GPU enjoys a higher throughput in half precision and `int8` precision, thanks to Tensor Cores usage.

Because the T4 and the P4 GPU have the same number of CUDA cores, we measure similar arithmetic throughput in matrix multiplication on the two boards, in double and single precision. Note that double-precision performance is hampered by the small number of native FP64 cores available (only two per SM), as both architectures are optimized for inference, where lower precision is more frequently employed.




Figure 4.3: Clock frequency observed on a T4 GPU while continuously computing cuBLAS matrix multiplication. The application clock frequency is set to 1,590 MHz.

4.5 Performance throttling

Most GPUs include forms of clock throttling and/or power-state throttling to prevent exceeding either the power or thermal envelopes if the workload is particularly demanding or the heat dissipation is insufficient.

Our experiments show that the small form-factor T4 and P4 boards, designed for inference applications, achieve a significantly higher frequency-per-Watt rating than their full-size counterparts. At the same time, they are more prone to clock throttling than their full-size counterparts (K80, P100, V100, and M60) because of

- their smaller size, which limits their heat sinks' heat transfer rate, and
- their maximum power limits set by the manufacturer, which is significantly lower (70W) on low-power, small form-factor boards than on full-size boards (250W).

Experimental setup: All GPU specimens we examined adopt passive cooling. Our K80, P100, V100 and M60 experiments ran on Dell PowerEdge C4130 servers, which




Figure 4.4: Temperature and clock frequency of the T4 card when computing a `cublasSgemm` repeatedly. The application clock frequency is set to 1,590 MHz.

are Tesla-qualified. Our T4 and P4 experiments ran on HPE Proliant DL360 Gen9 servers. This server model does not appear in NVidia’s Tesla-qualified server catalog. Power and thermal performance of a GPU also depend on the server that hosts it, and could be suboptimal on a non-qualified server. The server generation immediately following the one we employed (HPE Proliant DL360 Gen10) is Tesla-qualified, but we were unable to arrange for an upgrade before the publication of this manuscript.

In our experiments, we were able to trigger clock throttling on the T4 GPU consistently, using benchmarks based on cuBLAS matrix multiplication kernels `cublas<t>gemm`. On the T4 GPU, clock throttling triggers for two reasons:

- **power-limit throttling:** instantaneous power exceeds the power limit set by the manufacturer (70W on the T4 GPU);
- **thermal throttling:** the GPU reaches its maximum operating temperature (85°C on the T4 card).

Compared to power limit throttling, thermal throttling causes a more severe clock frequency reduction.




Figure 4.5: Clock frequency normalized to thermal design power (TDP) of all considered GPUs when computing an identical `cublas<v>sgemm` function on 1024×1024 matrices repeatedly. On every GPU, we set the application clock frequency to its maximum supported value.

4.5.1 Power-limit throttling

On the T4 and P4 GPUs, we saw power-limit throttling trigger very early in our cuBLAS-based matrix multiplication experiments. On the other hand, the V100, P100, M60 and K80 GPUs barely experienced any power-limit throttling, due to the larger margin between actual power consumption and its limit.

To confirm the cause of throttling, we designed an experiment that invokes `cuBLAS<v>gemm` kernels with input matrices of growing size. We observed the T4 GPU exceeded its power limit more and more frequently, and lower its clock rates more and more, with growing input sizes. The reduced clock frequency eventually hurts overall arithmetic throughput. See Figure 4.3.

In the experiment, we set the application clock for graphics on the T4 card to 1,590 MHz, and prevent GPU temperatures from exceeding the maximum operating temperature of the T4 GPU. We record the clock frequency of the T4 card while computing `cublas<v>gemm` in half precision.

4.5.2 Thermal throttling

We characterized thermal throttling with a benchmark that repeatedly launches a `cublas<t>gemm` kernel on a large matrix. We observed that below 85 degrees C (the maximum operating temperature), power limit throttling causes the T4 GPU to reduce its graphics clock with the growth of temperature. As soon as the temperature reaches 85 degrees C, thermal throttling triggers in addition to power-limit throttling, causing a more dramatic clock frequency step-down, depicted in Figure 4.4.

4.5.3 Power-limit throttling across GPU devices

We compared the power-limit throttling behavior of the different GPUs, by recording graphics clock over time while all cards computed endless repetitions of the same `cublasSgemm` kernel on 1024×1024 input matrices.

We noticed substantial differences between low-power GPUs (e.g., T4 and P4) and the full form-factor GPUs (K80, M60, P100, V100). We observe clock throttling only on the T4 and the P4 GPUs. Both cards are only able to run at their highest supported clock frequency for a few seconds at the very beginning of the experiment. As temperatures increased, clock throttling intervened and clock frequency decreased (Figure 4.5).

On full-height, full-length GPUs, we could not raise power consumption enough to approach the limits and trigger throttling.

Experimental setup: in all experiments, we set all graphics clocks to the highest supported value for each device. We turned off the AutoBoost features wherever available. We also ensured that only power-limit throttling was active.

Chapter 5

Conclusions

We refreshed our microbenchmark-based architectural discovery study, updating it for the Turing architecture. We revealed Turing’s architectural details, and compared them with previous NVidia architectures.

We emphasize the comparison between the T4 and the P4 GPUs: both are low-power, small-form-factor boards that target inference applications. The T4 is based on the Turing architecture; the P4, its predecessor, is based on Pascal.

We find that Turing uses the same instruction encoding as Volta, but it extends Volta’s instruction set; it also introduces a new register type (uniform registers) and supports more operand types on Tensor Cores. The new instructions allow the `nvcc` compiler to render matrix math in fewer instructions on Turing than on Volta.

The T4 GPU also delivers a significantly higher arithmetic throughput than the P4 on reduced-precision operands.

Turing’s memory hierarchy is similar to that of Volta, with different sizes at certain cache levels. We provided an exhaustive examination of the differences. When compared in terms of instruction encoding, memory hierarchy, and behavior of their processing units, the Turing and Volta generations display continuity, and together represent a significant departure from the Kepler and older generations.

Turing continues a trend of growth in the scheduler-to-cores ratio, which grew from 1:48 in Kepler to 1:16 in Turing. This trend correlates with a growth in instruction throughput. With their newly introduced L0 instruction cache, Turing and Volta mitigate the penalty associated with their longer instructions. The improved L1 data cache offers lower latency and higher bandwidth. Their new replacement policy also reduces cache miss rates when not using shared memory, and the change from 4, single-ported register banks in Pascal to 2, dual-ported banks facilitates the prevention of bank conflicts.

Compared to the Pascal P4 GPU, the Turing T4 GPU provides higher bandwidth on L1 cache and global memory. The T4 GPU has higher arithmetic throughput for half-precision, `int8` and `int4` matrix multiplication thanks to its improved Tensor Cores. In single and double precision, the T4 and the P4 GPUs exhibit comparable performance because they contain the same number of cores and are clocked at similar frequencies.

Using our findings on the instruction set encoding, software designers can optimize their code at the binary level and even construct customized SASS assemblers able to target Turing and possibly generate more tightly scheduled code, that delivers higher performance. Thanks to the memory hierarchy information we disclose, developers can also optimize their code by selecting working sets that match the cache memories at every suitable level, thus reducing miss rates and improving overall performance.

Appendix

In this appendix, we provide the opcodes for common instructions, as encoded in Turing's instruction encoding and, for comparison, in Pascal's and Volta's encoding.

Instruction	Floating point instructions		
	Pascal	Volta	Turing
FADD	0101 1100 0101 1 0100 1100 0101 1 0011 1001 0101 1	10 0010 0001	10 0010 0001
FCHK	0101 1100 1000 1 0100 1100 1000 1 0011 1001 1000 1	011 0000 0010	011 0000 0010
FCMP	0101 1011 1010 0101 0011 1010 0100 1011 1010 0011 0111 1010	–	–
FFMA	0101 1001 1 0101 0001 1 0100 1001 1 0011 0011 1 0011 0010 1	10 0010 0011	10 0010 0011
FMNMX	0101 1100 0110 0 0100 1100 0110 0 0011 1001 0110 0 0011 1000 0110 0	010 0000 1001	010 0000 1001
FMUL	0101 1100 0110 1 0100 1100 0110 1 0011 1001 0110 1 0011 1000 0110 1	010 0010 0000	010 0010 0000
FSET	0101 1000 0100 1000 0011 0001	010 0000 1010	010 0000 1010
FSETP	0101 1011 1011 0100 1011 1011 0011 0111 1011 0011 0110 1011	010 0000 1011	010 0000 1011
FSWZADD	0101 0000 1111 1	0 1000 0010 0010	0 1000 0010 0010
MUFU	0101 0000 1000 0	011 0000 1000	011 0000 1000
RRO	0101 1100 1001 0 0100 1100 1001 0	–	

Floating point instructions (continued)			
Instruction	Pascal	Volta	Turing
DADD	0011 1001 1001 0		
	0011 1000 1001 0		
	0101 1100 0111 0	10 0010 1001	10 0010 1001
	0100 1100 0111 0		
	0011 1001 0111 0		
DFMA	0011 1000 0111 0		
	0101 1011 0111	10 0010 1011	10 0010 1011
	0101 0011 0111		
	0100 1011 0111		
	0011 0111 0111		
DMNMX	0011 0110 0111		
	0101 1100 0101 0	–	–
	0100 1100 0101 0		
	0011 1001 0101 0		
	0011 1000 0101 0		
DMUL	0101 1100 1000 0	010 0010 1000	010 0010 1000
	0100 1100 1000 0		
	0011 1001 1000 0		
	0011 1000 1000 0		
	0101 1001 0	–	–
DSET	0100 1001 0		
	0011 0011 0		
	0011 0010 0		
	0101 1011 1000	10 0010 1010	10 0010 1010
	0100 1011 1000		
DSETP	0011 0111 1000		
	0011 0110 1000		
	–	10 0011 0000	10 0011 0000
	–	10 0011 0001	10 0011 0001
	–	0 0010 0011 0110	0 0010 0011 0110
HADD2	–	010 0011 0010	010 0011 0010
HFMA2	–	10 0011 0100	10 0011 0100
HMMA	–	10 0011 0011	10 0011 0011
HMUL2	–	010 0000 1000	010 0000 1000
HSETP2	–		
HSET2	–		
FSEL	–		

Integer Instructions			
Instruction	Pascal	Volta	Turing
BFE	0101 1100 0000 0	–	–
	0100 1100 0000 0		
	0011 1001 0000 0		
	0011 1000 0000 0		
	0101 1011 1111 0	–	–
BFI	0101 0011 1111 0		
	0100 1011 1111 0		
	0011 0111 1111 0		
	0011 0110 1111 0		
	0101 1100 0011 0	011 0000 0000	011 0000 0000
FLO	0100 1100 0011 0		
	0011 1001 0011 0		
	0011 1000 0011 0		

Integer Instructions (continued)			
Instruction	Pascal	Volta	Turing
IADD	0101 1100 0001 0 0100 1100 0001 0 0101 1100 0001 0 0101 1101 0001 0	–	–
IADD3	0101 1100 1100 0100 1100 1100 0011 1001 1100 0011 1000 1100	010 0001 0000	010 0001 0000
ICMP	0101 1011 0100 0101 0011 0100 0100 1011 0100 0011 0111 0100 0011 0110 0100	–	–
IMAD	0101 1010 0 0101 0010 0 0100 1010 0 0011 0100 0	10 0010 0100 10 0010 0101	10 0010 0100 10 0010 0101
IMADSP	0101 1010 1 0101 0010 1 0100 1010 1 0011 0101 1 0011 0100 1	–	–
IMNMX	0101 1100 0010 0 0100 1100 0010 0 0011 1001 0010 0 0011 1000 0010 0	–	–
IMUL	0011 1000 0011 1 0100 1100 0011 1 0011 1001 0011 1 0011 1000 0011 1	?	?
ISCADD	0101 1100 0001 1 0100 1100 0001 1 0011 1001 0001 1 0011 1000 0001 1	–	–
ISSET	0101 1011 0101 0100 1011 0101 0011 0111 0101 0011 0110 0101	–	–
ISSETP	0011 0111 0110 0100 1011 0110 0011 0111 0110 0011 0110 0110	010 0000 1100	010 0000 1100
LEA	0101 1011 1101 0 0101 1011 1101 1 0100 1011 1101 0 0011 0111 1101 0 0011 0110 1101 0 0001 1000	010 0001 0001	010 0001 0001
LOP3	0011 11 0101 1011 1110 0 0000 001	010 0001 0010	010 0001 0010
LOP	0101 1100 0100 0 0100 1100 0100 0		–

Integer Instructions (continued)			
Instruction	Pascal	Volta	Turing
POPC	0011 1001 0100 0	011 0000 1001	011 0000 1001
	0011 1000 0100 0		
	0101 1100 0000 1		
	0100 1100 0000 1		
	0011 1001 0000 1		
SHF	0011 1000 0000 1	10 0001 1001	10 0001 1001
	0101 1011 1111 1		
	0011 0111 1111 1		
	0011 1000 1111 1		
	0011 1001 1111 1		
SHL	0011 0110 1111 1	*	*
	0101 1100 1111 1		
	0101 1100 0100 1		
	0011 1000 0100 1		
	0011 1001 0100 1		
SHR	0100 1100 0100 1	*	*
	0101 1100 0010 1		
	0011 1000 0010 1		
	0011 1001 0010 1		
	0100 1100 0010 1		
XMAD	0101 1011 00	—	—
	0100 111		
	0101 0001 0		
	0011 0111 00		
	0011 0110 00		
VABSDIFF	—	10 0001 0100	10 0001 0100
VABSDIFF4	—	10 0001 0101	10 0001 0101
BREV	—	011 0000 0001	011 0000 0001
IABS	—	010 0001 0011	010 0001 0011
IDP	—	010 0010 0110	010 0010 0110
QSPC	—	0 0011 1010 1010	0 0011 1010 1010
BMSK	—	010 0001 1011	010 0001 1011

Conversion Instructions			
Instruction	Pascal	Volta	Turing
MOV	0101 1100 1001 1	010 0000 0010	010 0000 0010
	0100 1100 1001 1		
	0011 1001 1001 1		
	0011 1000 1001 1		
PRMT	0101 1011 1100	10 0001 0110	10 0001 0110
	0101 0011 1100		
	0100 1011 1100		
	0011 0111 1100		
SEL	0011 0110 1100	010 0000 0111	010 0000 0111
	0101 1100 1010 0		
	0011 1000 1010 0		
	0011 1001 1010 0		
SHFL	0100 1100 1010 0	0 1001 1000 1001	0 1001 1000 1001
	1110 1111 0001 0		
	0101 0000 1001 1		
	0101 0000 1010 0		
CSET	—	—	—
CSETP	—	—	—

Conversion Instructions (continued)			
Instruction	Pascal	Volta	Turing
PSET	0101 0000 1000 1	–	
PSETP	0101 0000 1001 0	–	
P2R	0101 1100 1110 1	010 0000 0011	010 0000 0011
	0100 1100 1110 1		
	0011 1001 1110 1		
	0011 1000 1110 1		
R2P	0101 1100 1111 0	010 0000 0100	010 0000 0100
	0100 1100 1111 0		
	0011 1001 1111 0		
	0011 1000 1111 0		
GETLMEMBASE	–	0 0011 1100 0000	0 0011 1100 0000

Load/Store Instructions			
Instruction	Pascal	Volta	Turing
LD	100	0 1001 1000 0000	0 1001 1000 0000
LDC	1110 1111 1001 0	0 1011 1000 0010	0 1011 1000 0010
LDG	1110 1110 1101 0	0 0011 1000 0001	0 0011 1000 0001
LDL	1110 1111 0100 0	0 1001 1000 0011	0 1001 1000 0011
LDS	1110 1111 0100 1	0 1001 1000 0100	0 1001 1000 0100
ST	101	0 0011 1000 0101	0 0011 1000 0101
STG	1110 1110 1101 1	0 0011 1000 0110	0 0011 1000 0110
STL	1110 1111 0101 0	0 0011 1000 0111	0 0011 1000 0111
STS	1110 1111 0101 1	0 0011 1000 1000	0 0011 1000 1000
ATOM	1110 1101	0 0011 1000 1010	0 0011 1000 1010
	1110 1110 011	0 0011 1000 1011	0 0011 1000 1011
	1110 1110 1111		
ATOMS	1110 1100	0 0011 1000 1100	0 0011 1000 1100
	1110 1110 00	0 0011 1000 1101	0 0011 1000 1101
	1110 1110 010		
ATOMG	–	0 0011 1010 1000	0 0011 1010 1000
		0 0011 1010 1001	0 0011 1010 1001
RED	1110 1011 1111 1	0 1001 1000 1110	0 1001 1000 1110
CCTL	1110 1111 0111	0 1001 1000 1111	0 1001 1000 1111
MEMBAR	1110 1111 1001 1	0 1001 1001 0010	0 1001 1001 0010
ERRBAR	–	0 1001 1010 1011	0 1001 1010 1011
CCTL	–		0 1001 1000 1111
CCTLT	1110 1011 1111 0		
CCTLL		0 1001 1001 0000	
MATCH	–	0 0011 1010 0001	0 0011 1010 0001

Control Instructions			
Instruction	Pascal	Volta	Turing
BRA	1110 0010 0100	0 1001 0100 0111	0 1001 0100 0111
BRX	1110 0010 0101	0 1001 0100 1001	0 1001 0100 1001
JMP	1110 0010 0001	0 1001 0100 1010	0 1001 0100 1010
JMX	1110 0010 0000	0 1001 0100 1100	0 1001 0100 1100
SSY	1110 0010 1001		
SYNC	1111 0000 1111 1	–	

Control Instructions (continued)			
Instruction	Pascal	Volta	Turing
BSYNC	–	0 1001 0100 0001	0 1001 0100 0001
WARPSYNC	–	011 0100 1000	011 0100 1000
CAL	1110 0010 0110	–	–
CALL	–	011 0100 0011	011 0100 0011
		0 1001 0100 0100	0 1001 0100 0100
JCAL	1110 0010 0010	–	–
PRET	1110 0010 0111	–	–
RET	1110 0011 0010	0 1001 0101 0000	0 1001 0101 0000
BRK	1110 0011 0100	–	–
PBK	1110 0010 1010	–	–
CONT	1110 0011 0101	–	–
PCNT	1110 0010 1011	–	–
EXIT	1110 0011 0000	0 1001 0100 1101	0 1001 0100 1101
PEXIT	1110 0010 0011	–	–
BPT	1110 0011 1010	–	–
BMOV	–	0 0011 0101 0101	0 0011 0101 0101
		011 0101 0110	011 0101 0110
		011 0101 0111	011 0101 0111
YIELD	–	0 1001 0100 0110	0 1001 0100 0110
RTT	–	0 1001 0100 1111	0 1001 0100 1111
KILL	–	0 1001 0101 1011	0 1001 0101 1011
RPCMOV	–	011 0101 0010	011 0101 0010
		0 0011 0101 0011	0 0011 0101 0011
		–	0 1001 0101 0100
IDE	–	0 1001 0101 0001	0 1001 0101 0001
PMTRIG	–	0 1000 0000 0001	0 1000 0000 0001
BREAK	–	0 1001 0100 0010	0 1001 0100 0010
BSSY	–	0 1001 0100 0101	0 1001 0100 0101
NANOSLEEP	–	011 0101 1101	011 0101 1101
NANOTRAP	–	011 0101 1010	011 0101 1010

Other Instructions			
Instruction	Pascal	Volta	Turing
NOP	0101 0000 1011 0	0 1001 0001 1000	0 1001 0001 1000
CS2R	0101 0000 1100 1	0 1000 0000 0101	0 1000 0000 0101
S2R	1111 0000 1100 1	0 1001 0001 1001	0 1001 0001 1001
B2R	1111 0000 1011 1	0 0011 0001 1100	0 0011 0001 1100
BAR	1110 0010 0100	011 0001 1101	011 0001 1101
R2B	1111 0000 1100 0	0 0011 0001 1110	0 0011 0001 1110
VOTE	0101 0000 1101 1	0 1000 0000 0110	0 1000 0000 0110
	0101 0000 1110 0	–	–
TMML	–	0 1011 0110 1001	0 1011 0110 1001
TXD	–	0 1011 0110 1100	0 1011 0110 1100
SGXT	–	010 0001 1010	010 0001 1010
AL2P	–	–	0 1001 0010 0000
CSMTEST	–	0 1000 0000 1101	0 1000 0000 1101
DEPBAR	–	0 1001 0001 1010	0 1001 0001 1010
IPA	–	–	0 0011 0010 0110
ISBERD	–	–	0 1001 0010 0011
LEPC	–	0 0011 0100 1110	0 0011 0100 1110
OUT	–	–	0 0011 0010 0100

Other Instructions (continued)			
Name	Pascal	Volta	Turing
PIXLD	–	–	0 1001 0010 0101
PLOP3	–	0 1000 0001 1100	0 1000 0001 1100
SETCTAID	–	0 0011 0001 1111	0 0011 0001 1111
SETLMEMBASE	–	0 0011 1100 0001	0 0011 1100 0001

Bibliography


- [1] Z. Jia, M. Maggioni, B. Staiger, and D. P. Scarpazza, “Dissecting the NVidia Volta GPU architecture via microbenchmarking,” in *2018 GPU Technology Conference*, 2018. [Online]. Available: <http://on-demand.gputechconf.com/gtc/2018/presentation/s8122-dissecting-the-volta-gpu-architecture-through-microbenchmarking.pdf>
- [2] —, “Dissecting the NVidia Volta GPU architecture via microbenchmarking,” 2018. [Online]. Available: <https://arxiv.org/abs/1804.06826>
- [3] *NVidia Turing GPU Architecture*. NVIDIA Corporation, 2018.
- [4] Y. Hou, “Asfermi,” 2011. [Online]. Available: <https://github.com/hyqneuron/asfermi>
- [5] C. Wang, Z. Jia, and K. Chen, “Tuning performance on Kepler GPUs: An introduction to Kepler assembler and its usage in CNN optimization,” in *GPU Technology Conference Presentation*, vol. 6173, 2015.
- [6] S. Gray, “maxas,” 2016. [Online]. Available: <https://github.com/NervanaSystems/maxas/wiki/Control-Codes>
- [7] H. Wong, M. M. Papadopoulou, M. Sadooghi-Alvandi, and A. Moshovos, “Demystifying GPU microarchitecture through microbenchmarking,” in *2010 IEEE International Symposium on Performance Analysis of Systems Software (ISPASS)*, March 2010, pp. 235–246.
- [8] X. Zhang, G. Tan, S. Xue, J. Li, K. Zhou, and M. Chen, “Understanding the GPU microarchitecture to achieve bare-metal performance tuning,” in *Proceedings of the 22nd ACM SIGPLAN Symposium on Principles and Practice of Parallel Programming*, ser. PPOPP ’17. New York, NY, USA: ACM, 2017, pp. 31–43.
- [9] X. Mei and X. Chu, “Dissecting GPU memory hierarchy through microbenchmarking,” *IEEE Transactions on Parallel and Distributed Systems*, vol. 28, no. 1, pp. 72–86, Jan 2017.

- [10] *NVidia Tesla V100 GPU Architecture, The World's Most Advanced Data Center GPU.* NVIDIA Corporation, 2017.
- [11] "NVidia cuobjdump and nvdisasm," 2016. [Online]. Available: <https://docs.nvidia.com/cuda/cuda-binary-utilities/>

Contents

1	Low-level details make a difference	5
2	How Turing encodes instructions	9
2.1	Control information	10
2.2	Processing Blocks and Schedulers	12
2.3	Instruction word format	13
3	Memory hierarchy	15
3.1	L1 data cache	18
3.2	Unified L2 cache	21
3.3	Instruction cache hierarchy	25
3.4	Constant memory hierarchy	28
3.5	Registers	30
3.6	Shared memory	33
3.7	Global memory	35
3.8	TLBs	36
4	Instruction latency and throughput	39
4.1	Native instructions	39
4.2	Atomic operations	40
4.3	New Tensor Core instructions	43
4.4	Arithmetic performance	44
4.5	Performance throttling	46
5	Conclusions	51
	Bibliography	61
	Contents	63

The Authors



Zhe Jia is a senior R&D engineer with the High-Performance Computing group at Citadel. Prior to this position, he was a senior R&D engineer with Alicloud, and a software engineering intern at Microsoft, Beijing. He received his B.S. degree in Physics and his M.S. degree in Meteorology from Peking University, China. His interests include the performance optimization of deep learning, numerical modeling, and atmospheric simulation workloads.



Marco Maggioni Marco is a senior R&D engineer with the High-Performance Computing group at Citadel, Chicago. He received his Ph.D. in Computer Science from the University of Illinois at Chicago, where he focused on sparse linear algebra and convex optimization on GPUs. His research established records for the fastest sparse matrix-vector multiplication GPU kernel in research and industry.



Jeffrey Smith is a senior software engineer with the Market-Making Trading Platform R&D team at Citadel Securities. He received his B.S. in Computer Science from Rose-Hulman Institute of Technology. His interests include hardware architectural analysis, latency-focused systems acceleration and design, and software framework creation for thread- and network-level task parallelization.



Daniele Paolo Scarpazza Daniele leads the High-Performance Computing group at Citadel, Chicago. Prior to this position, he was a Research Scientist with D. E. Shaw Research, a Research Staff Member with the IBM T. J. Watson Research Center, and a Post-Doc with the Pacific Northwest National Laboratory. He received his Ph.D. in Information Engineering from Politecnico di Milano, Italy. He is the co-recipient of a Gordon Bell prize. He focuses on quantitative performance analysis and optimization of diverse algorithms on parallel architectures.